

# A3 REFLECTION + INTRO TO MODEL-BASED

---

CSE 599 Prototyping Interactive Systems | Lecture 14 | May 1

**Jon Froehlich** • Jasper O'Leary (TA)

Spring 2019

Home

Announcements

Assignments

Discussions

Grades

People

Pages

Files

Syllabus

Outcomes

Quizzes

Modules

Conferences

Collaborations

Chat

Attendance

UW Libraries

# A3: Signal Processing + Machine Learning 1: Offline Gesture Recognizer

PublishedEdit⋮

## Overview

Imagine working for a new hardware startup designing new input controllers. You've been asked to prototype a new, custom input device with gesture recognition (using an accelerometer)--for example, to use the input controller as a paddle in tennis, as a "bat" in bowling or to recognize the "overhand throwing motion" in baseball. In this assignment you will build your own offline gesture recognizer to automatically recognize these gestures.

# A3 SHARE OUT!

While in the "real world" you would ultimately need to create a real-time gesture recognizer, for this assignment you will make an *offline* version in [Jupyter Notebook](#) (we strongly recommend the Anaconda Distribution). Specifically, you will make two recognizers:

- (i) a *shape-matching* recognizer such as via a Euclidean distance metric or Dynamic Time Warping and
- (ii) a *feature-based* (or *model-based*) recognizer using a [support-vector machine \(SVM\)](#) (recommended) or an alternative supervised learning approach of your choosing (e.g., an HMM).

Within Jupyter Notebook, we will use Python 3 and these amazing libraries [numpy](#), [scipy](#), [matplotlib](#), and [scikit-learn](#). Numpy and scipy provide numeric array handling and signal processing, matplotlib provides visualization, and scikit-learn is the de facto machine learning library in Python. You are welcome to use other libraries as well (e.g., [this DTW library](#)).

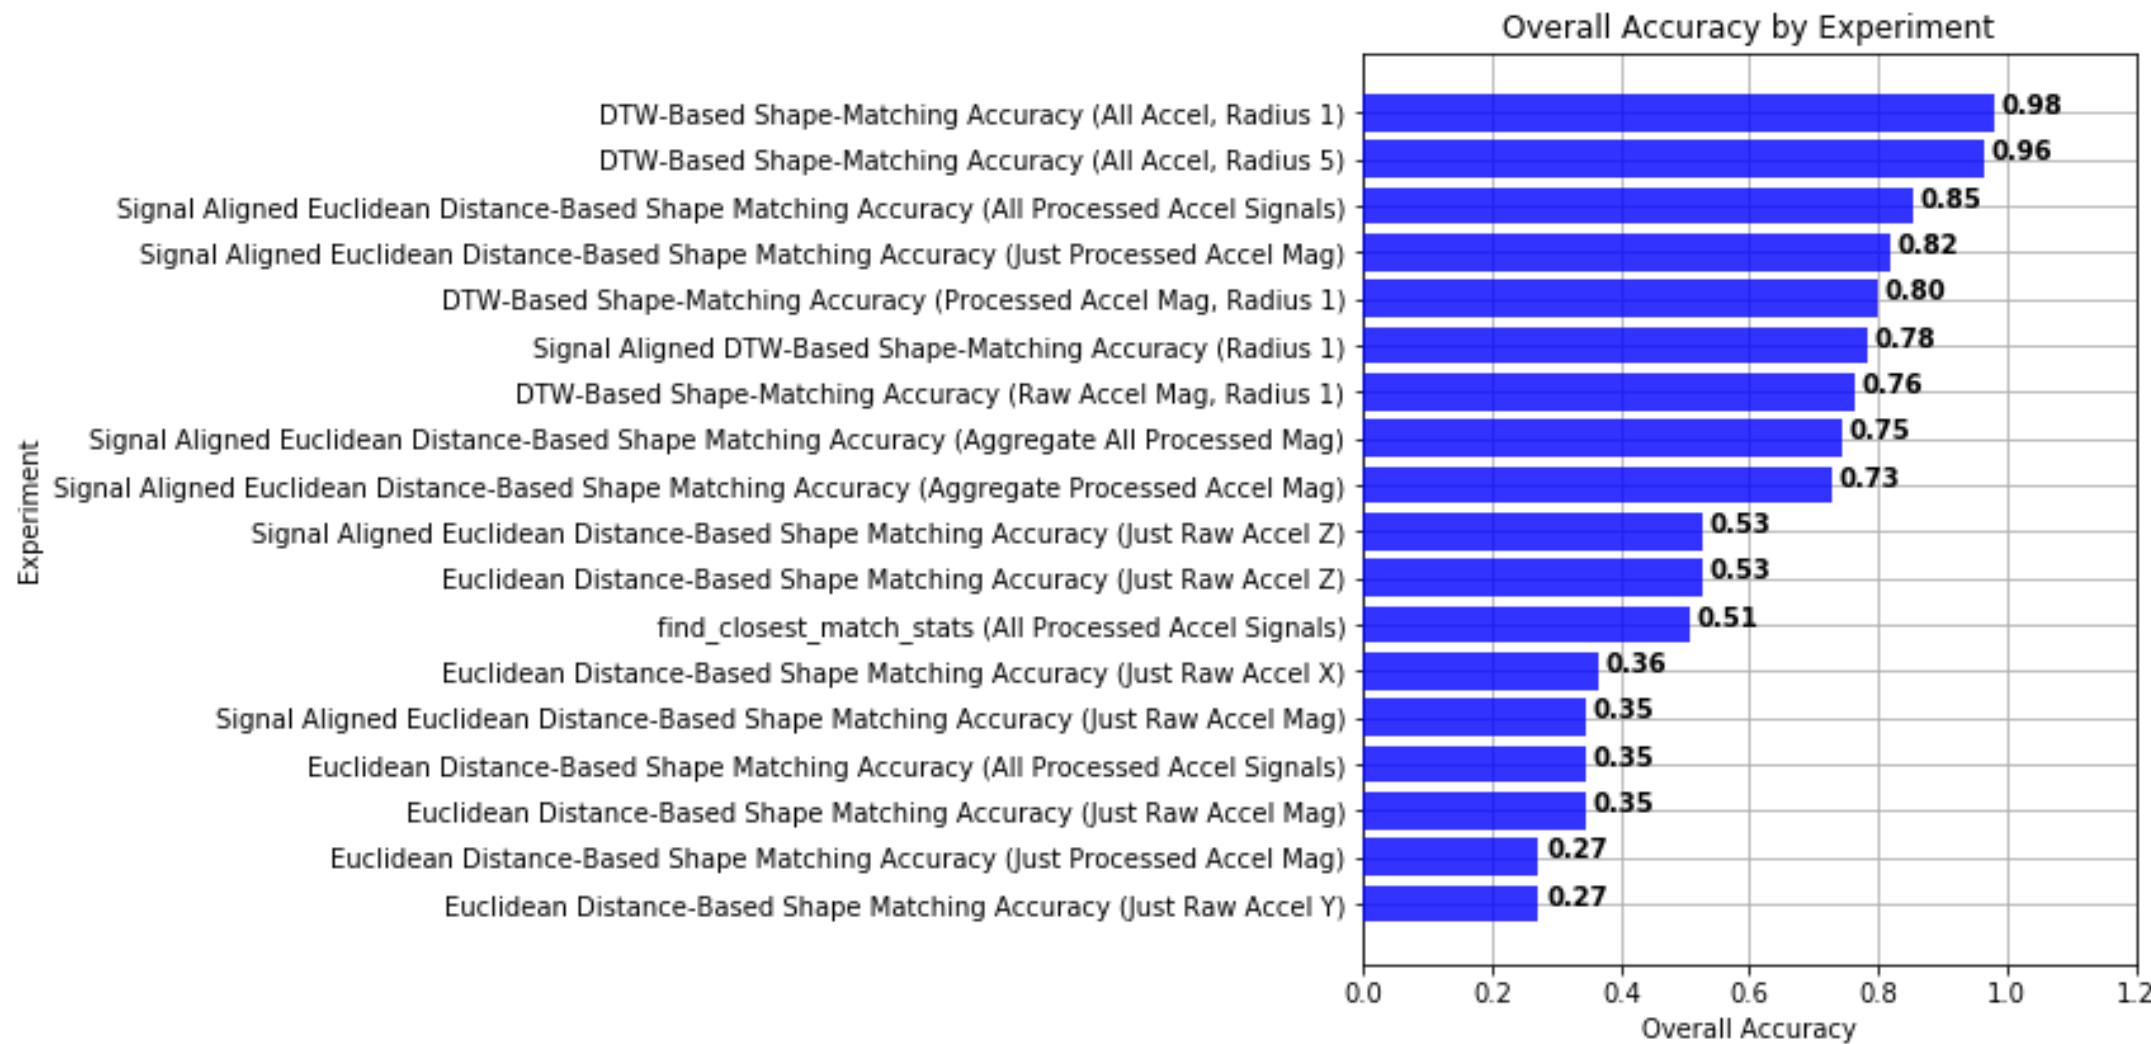
For your deliverables, you will turn in your Jupyter Notebook, your recorded gestures, and a slide deck

Related Items

SpeedGrader™

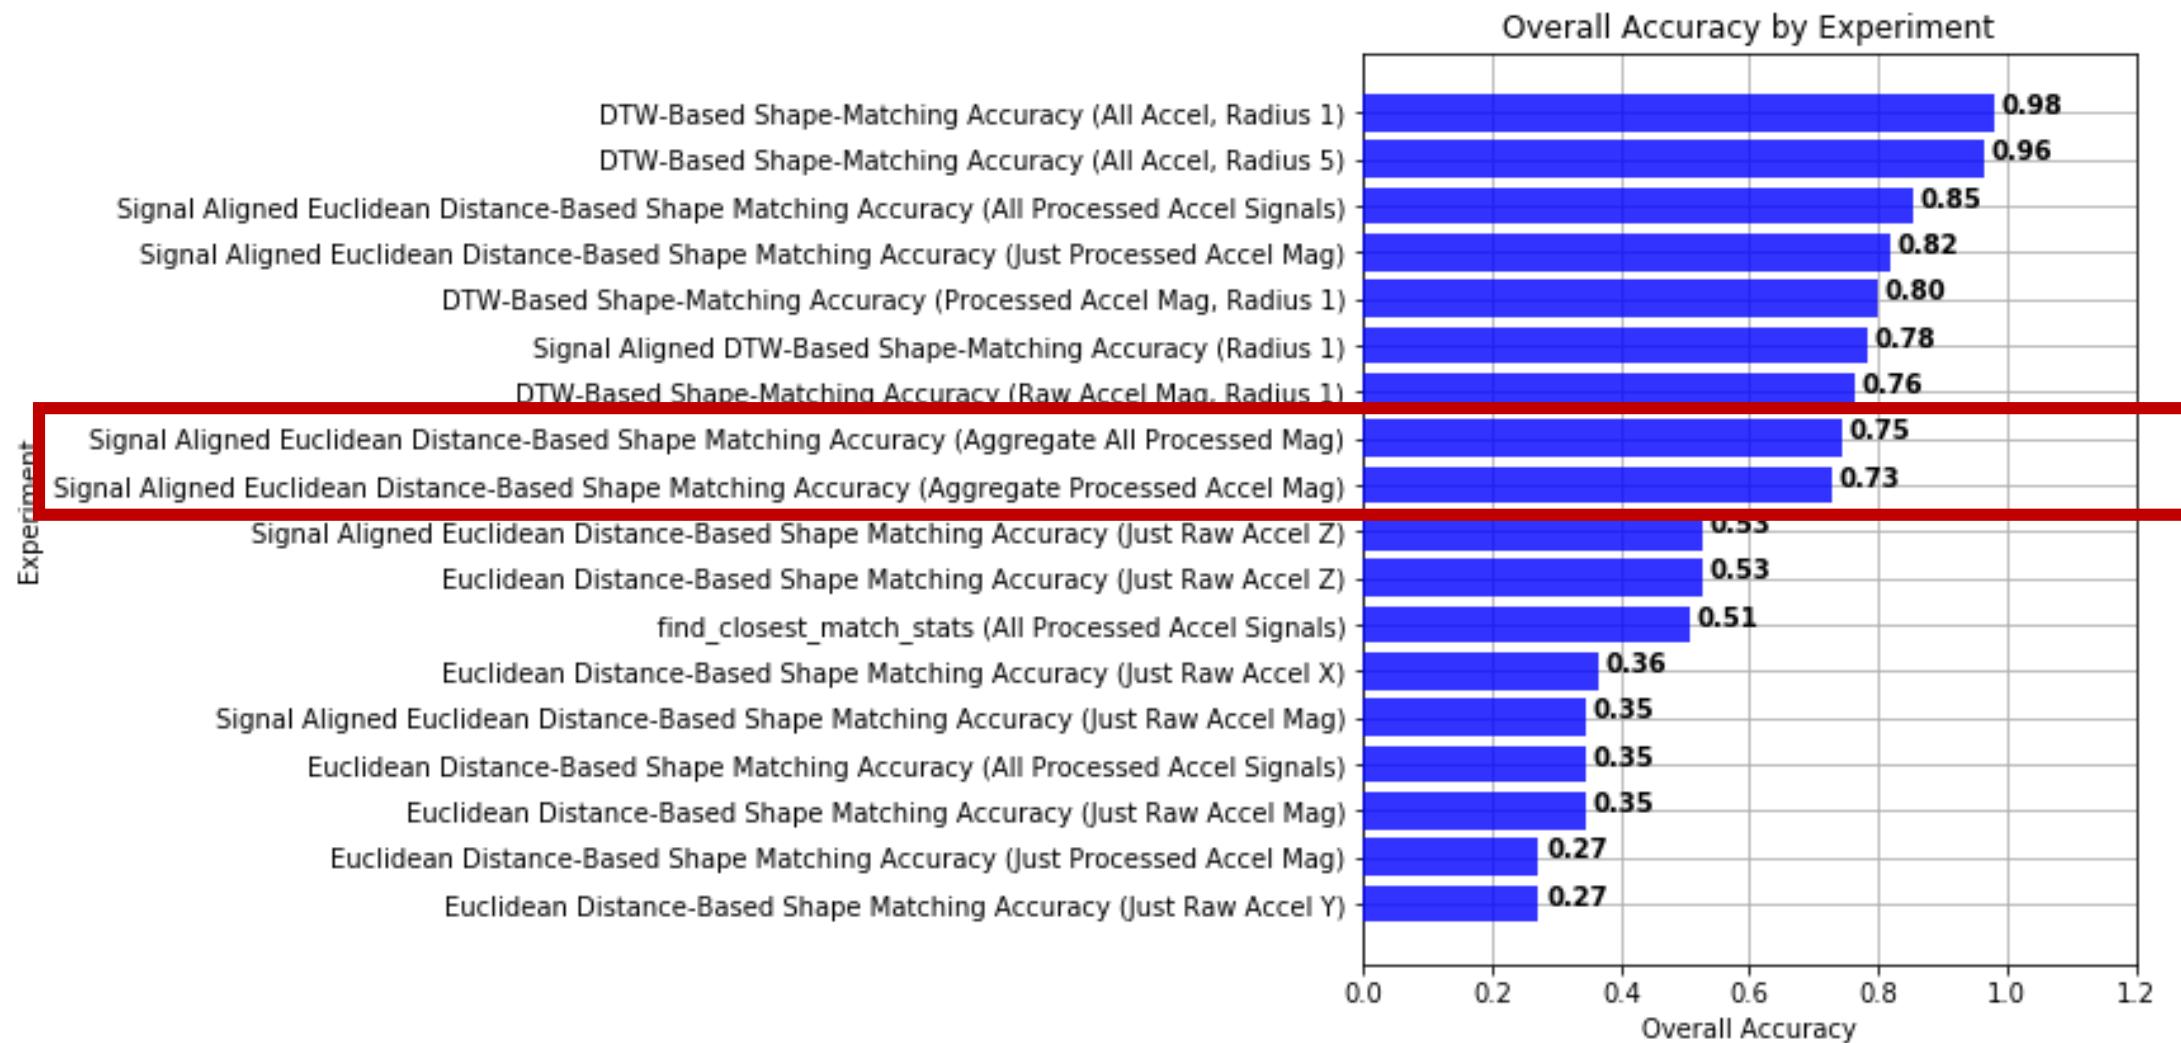
### A3: REFLECTION

# BEST PERFORMING ON MY “HARD” GESTURE SET

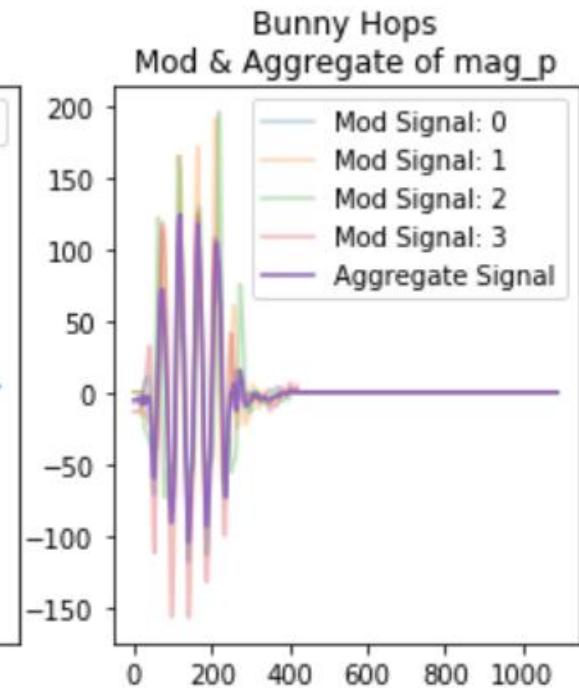
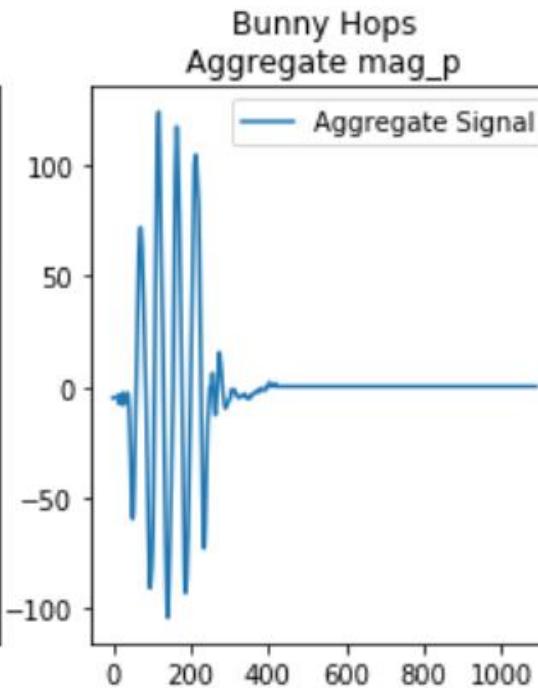
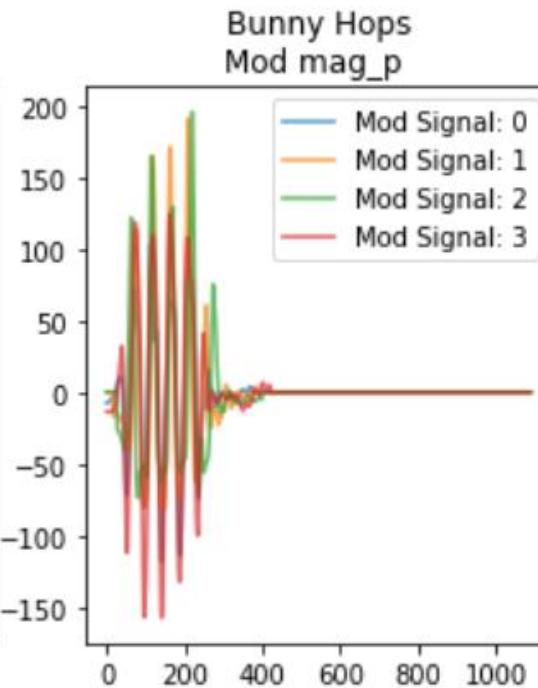
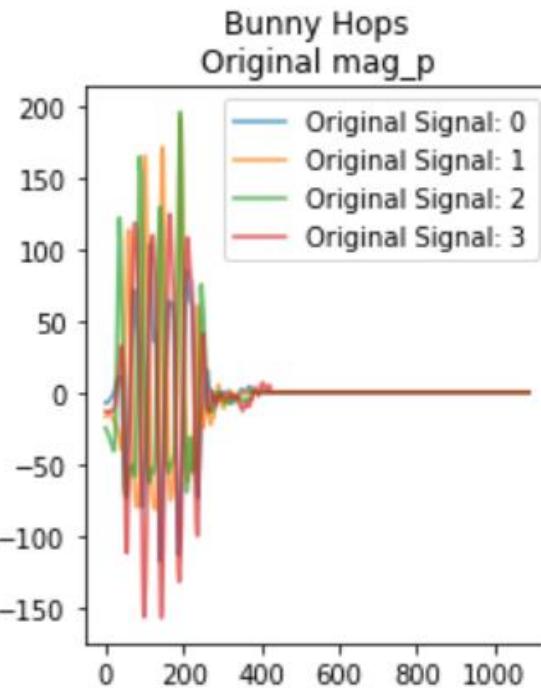


### A3: REFLECTION

# BEST PERFORMING ON MY “HARD” GESTURE SET

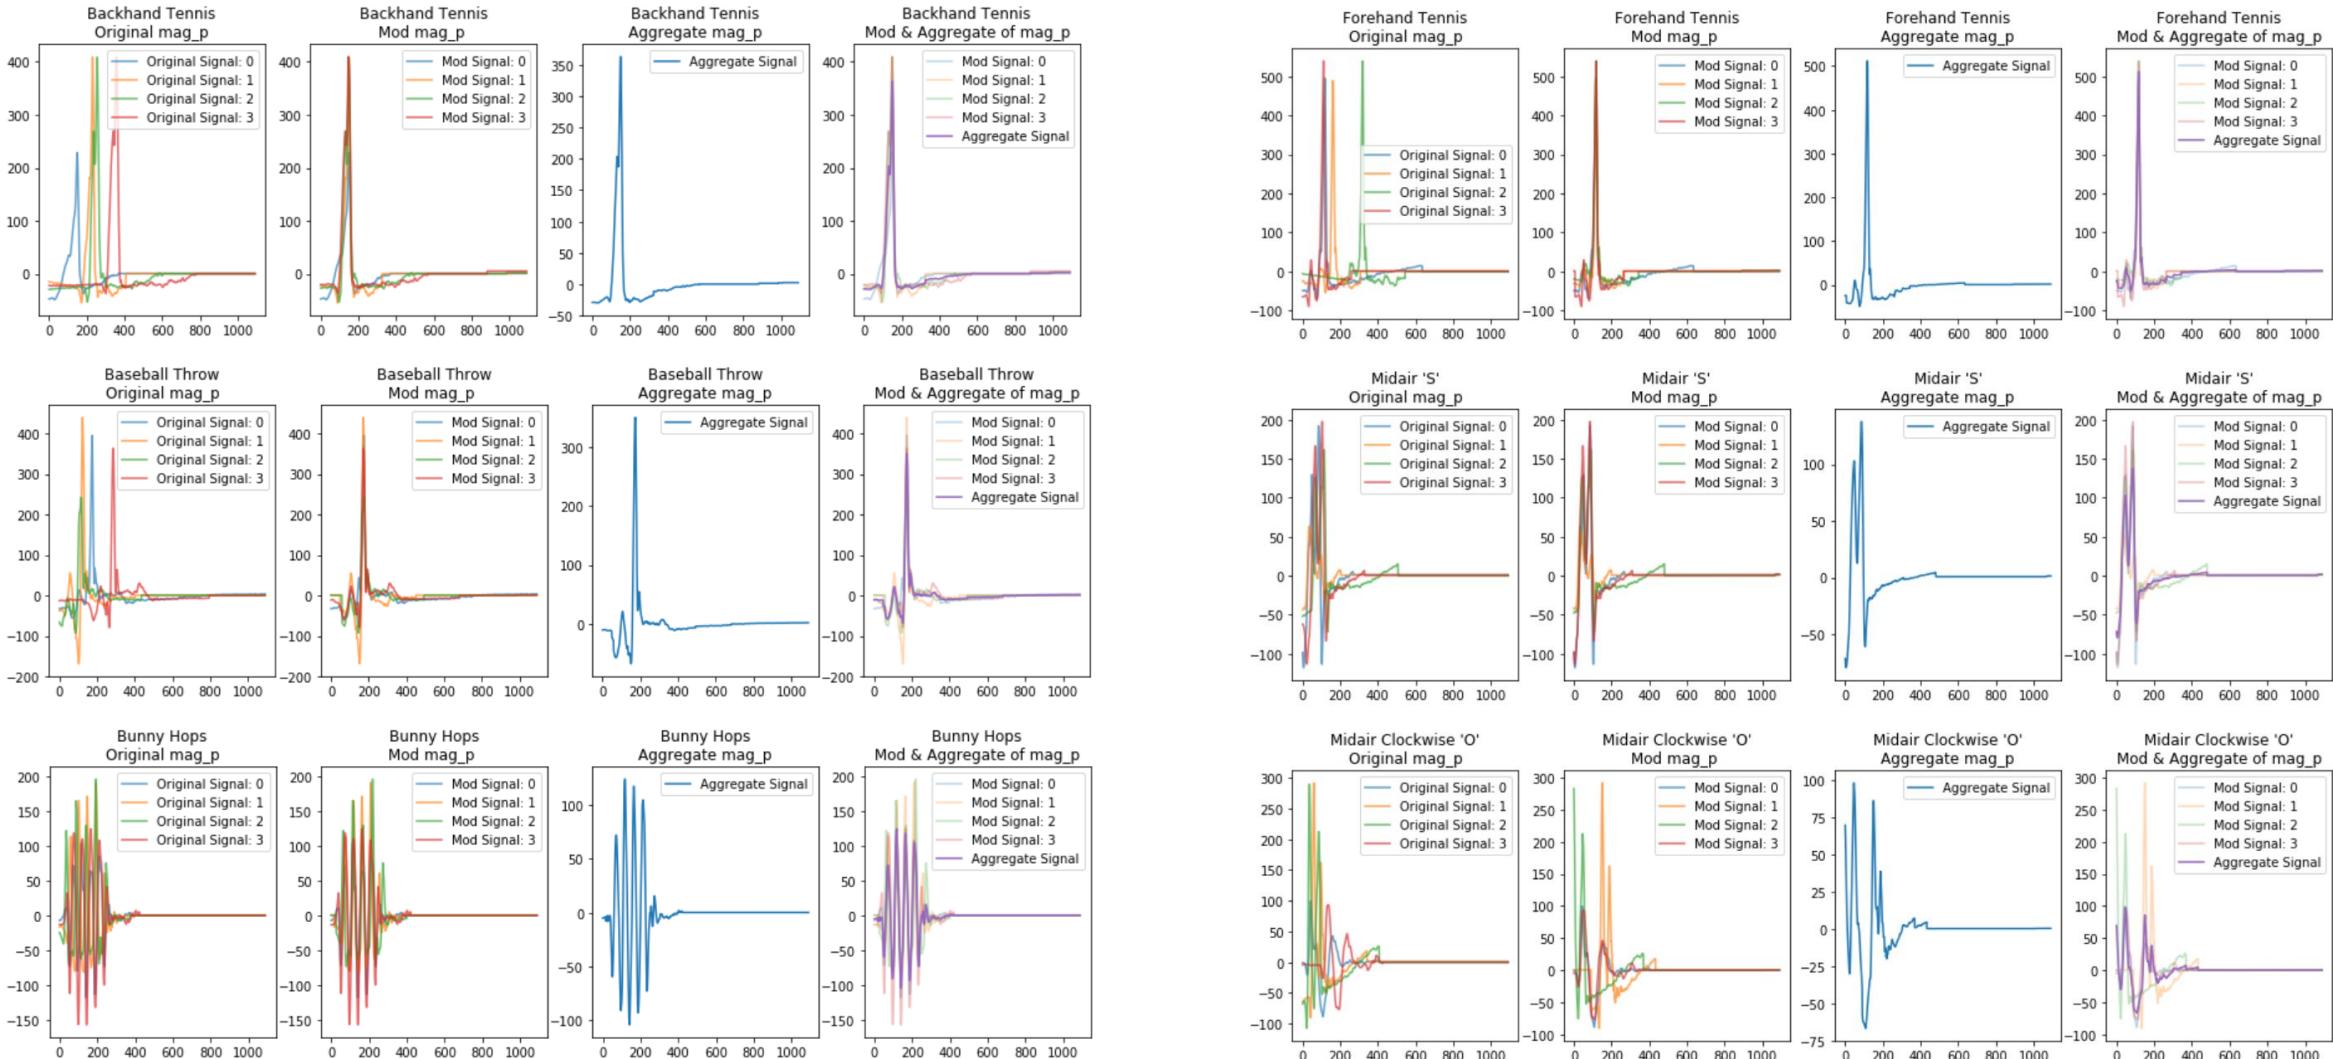


# CREATING AGGREGATE MODEL FROM TRAINING FOLDS

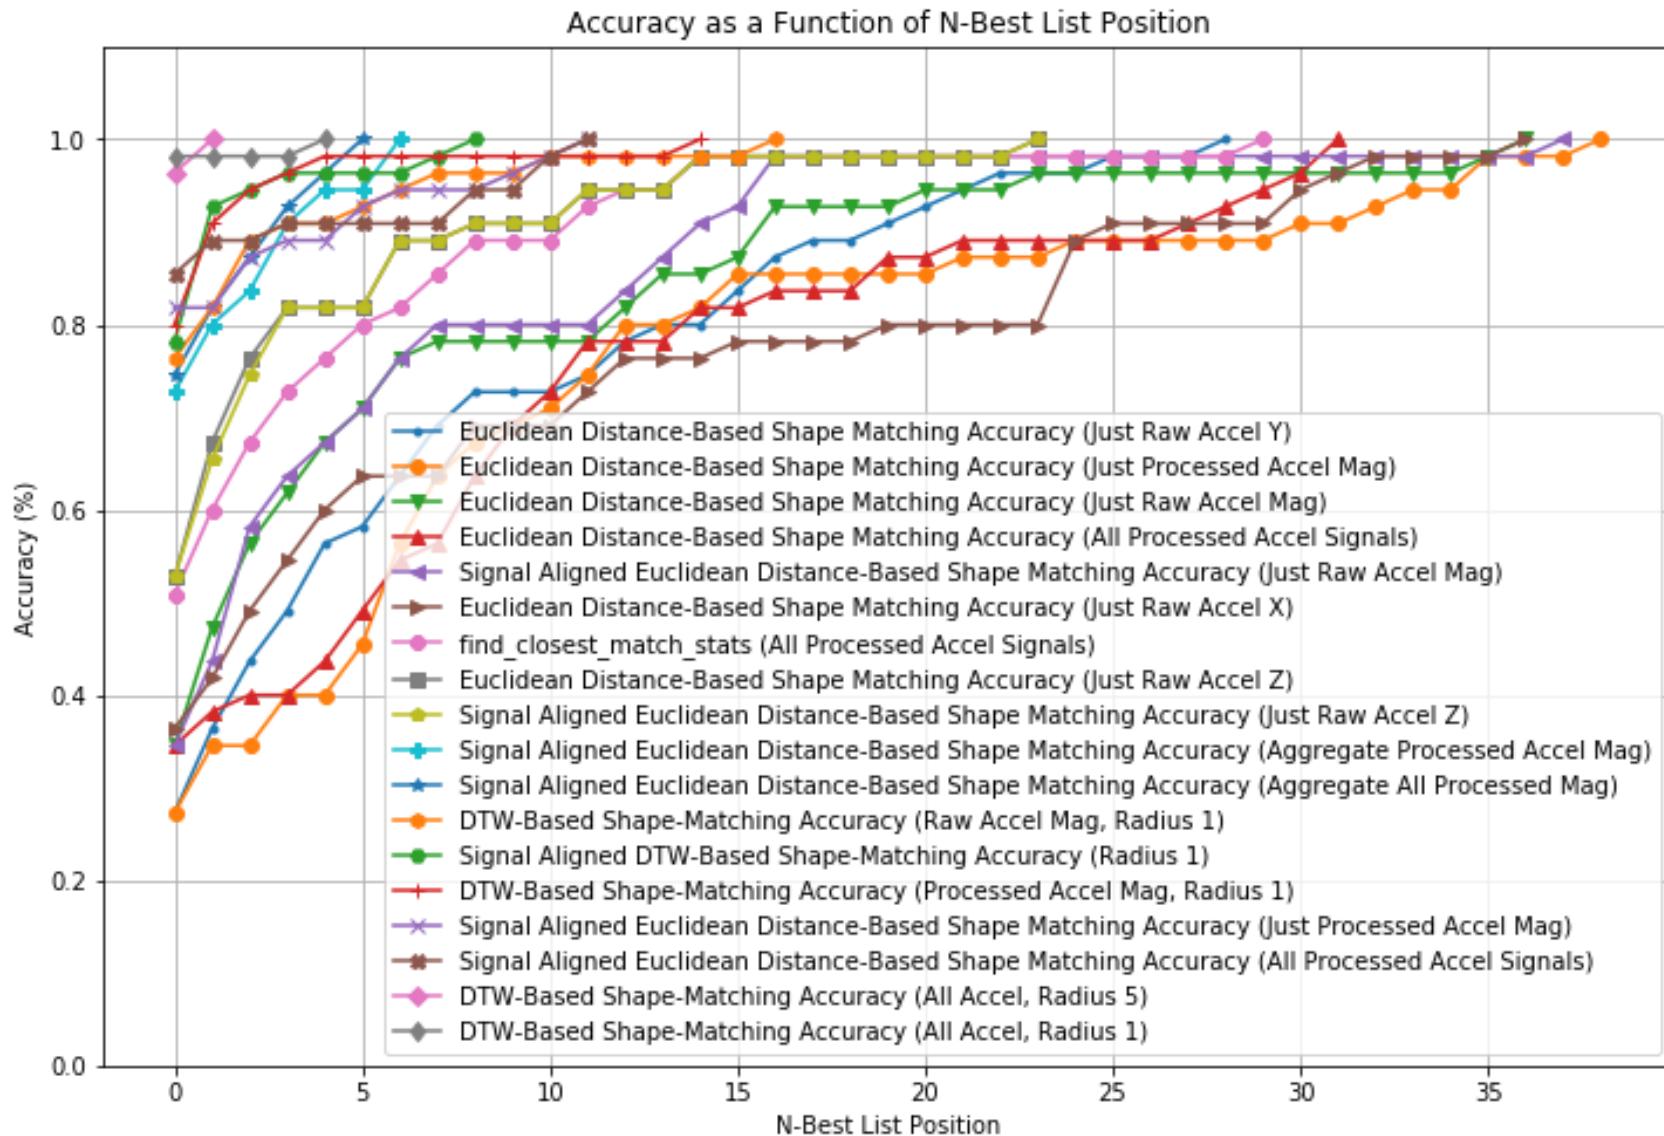


### A3: REFLECTION

# CREATING AGGREGATE MODEL FROM TRAINING FOLDS

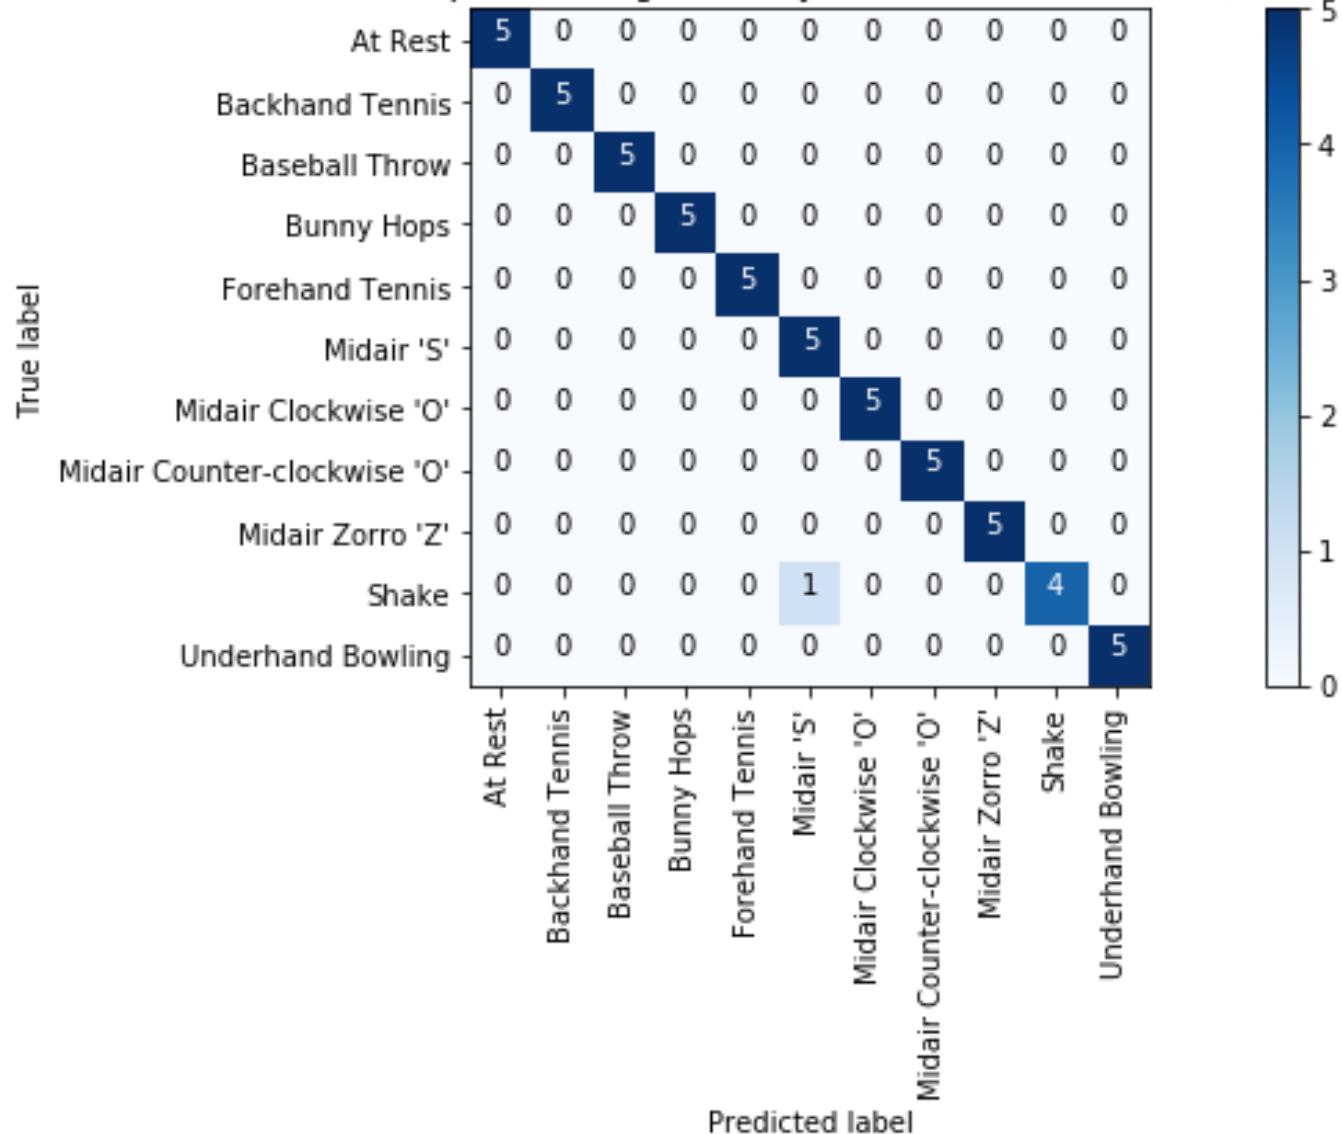


# BEST PERFORMING ON MY “HARD” GESTURE SET

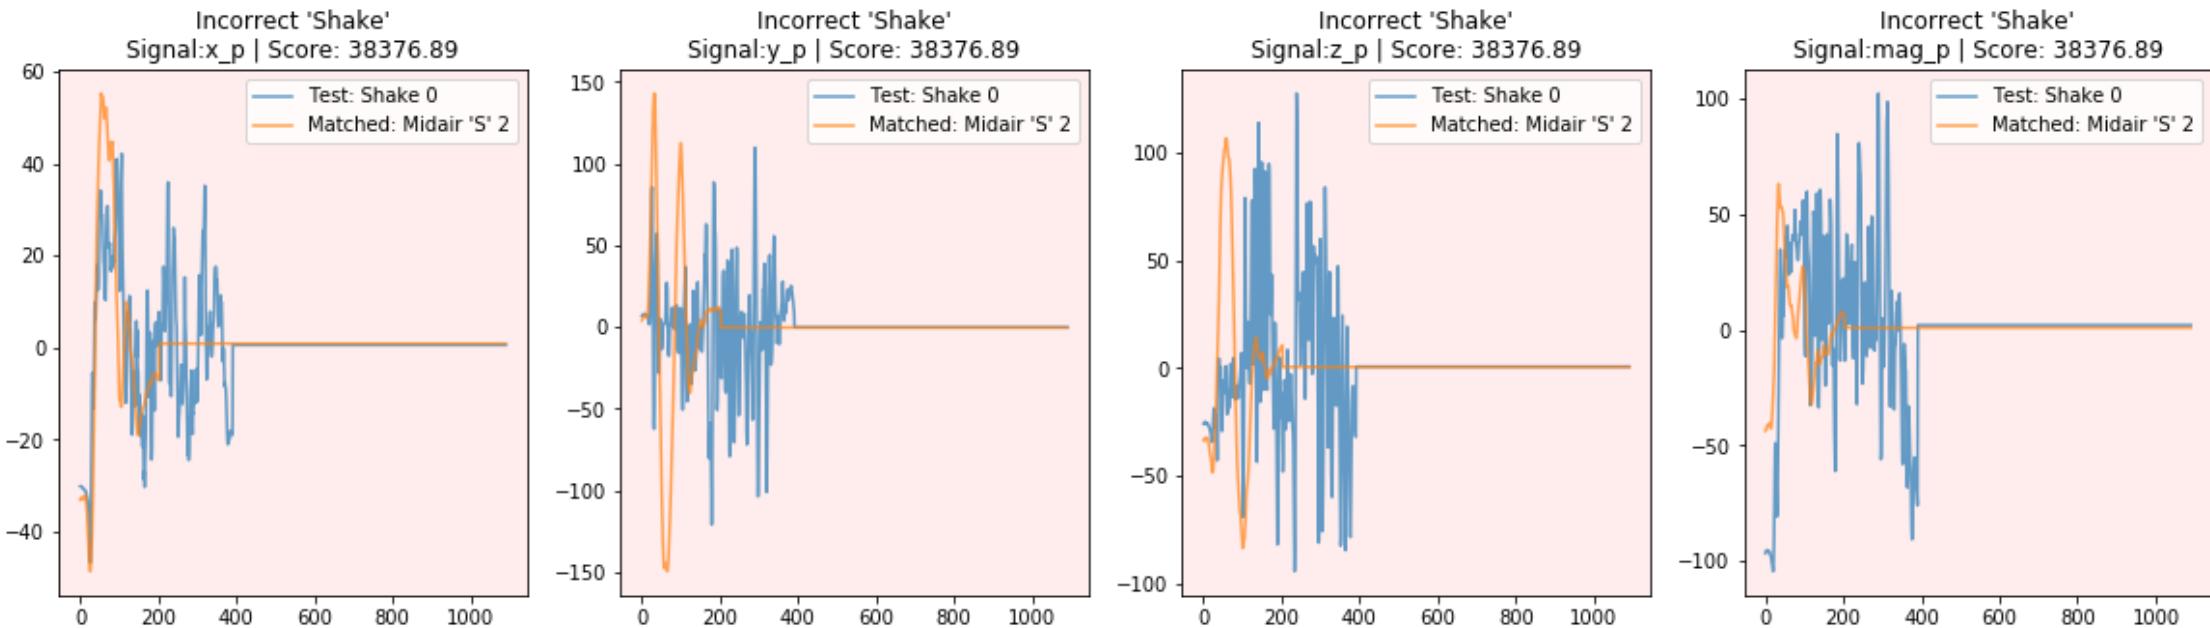


# CONFUSION MATRIX FOR CLASSIFYING ‘HARD’ GESTURE SET

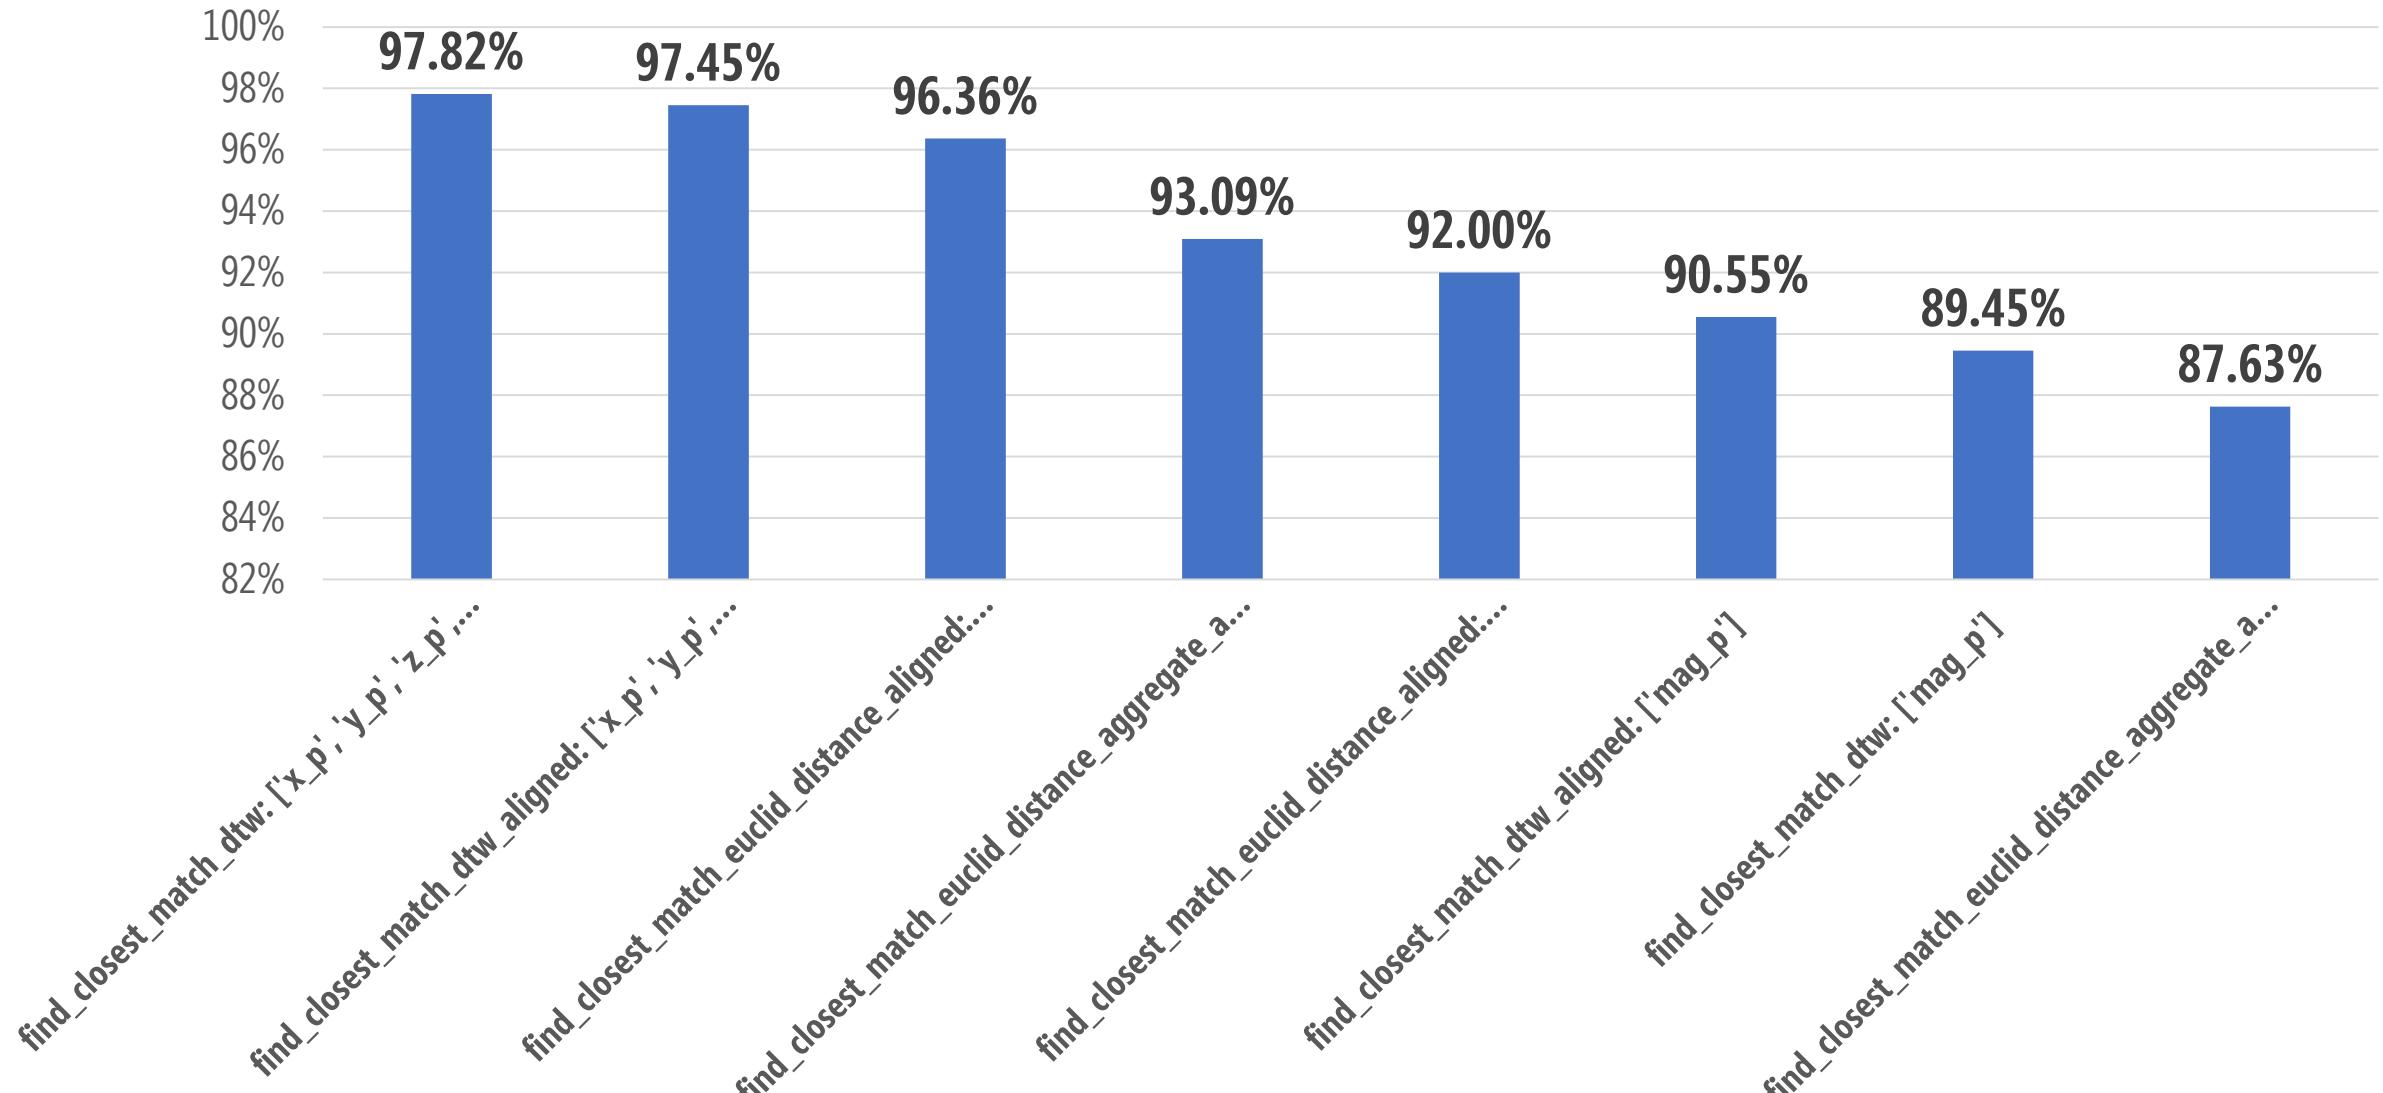
DTW-Based Shape-Matching Accuracy (All Accel, Radius 1): 54/55 (98.18%)



# VISUALIZING THE MISTAKE... CORRECTABLE?

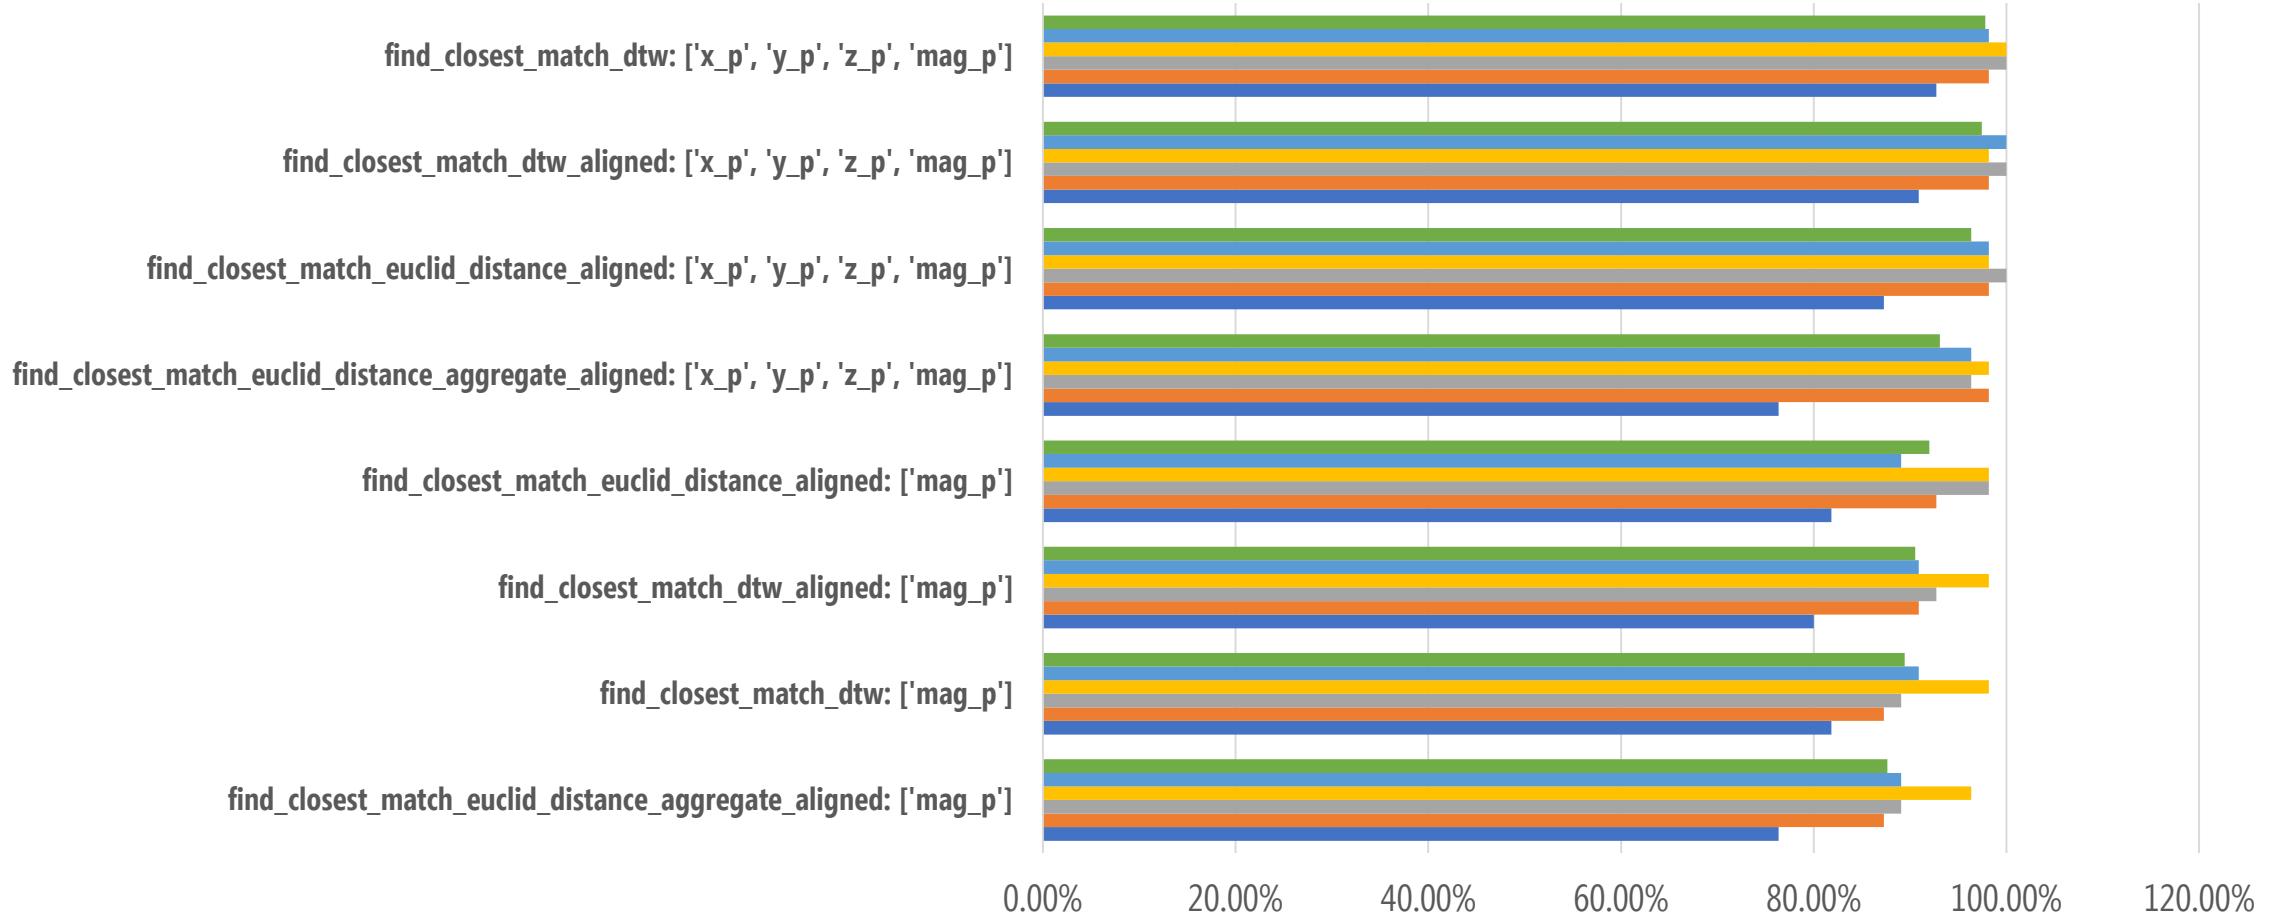


# OVERALL BEST ACROSS 3 STUDENTS + MY 2 GESTURE SETS



# OVERALL BEST ACROSS 3 STUDENTS + MY 2 GESTURE SETS

Overall    NicoleGestureLogs    JohnGestureLogs    JonGestureLogsEasy    AnnieGestureLogs    JonGestureLogsHarder



TODAY

# A3 REFLECTION + INTRO TO MODEL-BASED

**A3 reflection.** What worked? What didn't? Challenges? Solutions?

Quick intro to **model-based approaches**

Jupyter Notebook **exercises** with [SVMGestureRecognizer.ipynb](#)

Spring 2019

Home

Announcements

**Assignments**

Discussions

Grades

People

Pages

Files

Syllabus

Outcomes

Quizzes

Modules

Conferences

Collaborations

Chat

Attendance

UW Libraries

Add 4.0 Grade Scale

Panopto Recordings

Settings

# A4: Signal Processing + Machine Learning 2: Model-Based Gesture Recognizer

**Published****Edit**

...

Related Items

SpeedGrader™

## Overview

This assignment directly builds on A3: you will build a *feature-based* (or *model-based*) recognizer using a [support-vector machine \(SVM\)](#) (recommended) or an alternative supervised learning approach of your choosing (e.g., an HMM). You will also compare your results to the shape-matching algorithm from A3. At a minimum, you must test on my gesture set (11 gestures) and your gesture set (11 gestures). You need to upload your gesture set to this [Google Drive folder](#) as part of the assignment.

## INTRO TO A4

To get started, download and open this basic [SVM Gesture Classifier.ipynb](#). Like in A3, we provided you with some parsing code, data structures, and some initial visual explorations. In addition, we have also built the full training+test pipeline for training and testing the SVM using k-fold cross validation. Note: for the file handling code to work, you must have a folder called 'GestureLogs' in the same dir as this Notebook. Within 'GestureLogs', place folders of your gesture sets (at the very least, mine and yours but you could also download others from the [Google Drive](#))

You need to decide how to incorporate your A3 shape-matching approach into this notebook (e.g., you could make a copy of your A3 assignment notebook and then copy over relevant SVM code or you could copy over your A3 code into this notebook).

## Learning Goals

- Reinforce signal processing skills from A2 and complement them frequency signal analysis
- Introduce and learn feature extraction and feature classification concepts
- Introduce and learn notions of training + test and reinforce use of k-fold cross validation

## Parts

## A4: MODEL-BASED CLASSIFICATION

# MODEL-BASED RECOGNIZER

Drastically reduced scope of A4 to help you focus on your projects. No longer need to do event segmentation. I also added a full end-to-end pipeline to train and test your SVM.

Build a model-based (aka feature-based) recognizer using a Support Vector Machine

Compare your shape-matching results to this new approach

Upload your gesture set to the Google Drive

### The SVM

```
In [9]: # This is the simplest possible SVM using only a few features but gives you a sense of the overall approach
# Some nice resources:
# - A very simple classification example using scikit:
#   https://daumgartel.wordpress.com/2014/03/10/a-scikit-learn-example-in-10-lines/
# - A nice video overview of SVM: https://youtu.be/NIV0golbjSc
# - Official scikit Learn: http://scikit-learn.org/stable/modules/svm.html

from sklearn import svm
import itertools

# Returns a feature vectof for the given trial
def extract_features_example(trial):

    # Play around with features to extract and use in your model
    # Brainstorm features, visualize ideas, try them, and iterate
    features = []
    features.append(trial.accel_mag.max())
    features.append(np.std(trial.accel_mag))
    return features

numFolds = 5
selected_gesture_set = get_gesture_set_with_str("Hard")
numGestures = selected_gesture_set.get_num_gestures()
numTrialsTotal = selected_gesture_set.get_total_num_of_trials()

# Setting a seed here keeps producing the same folds set each time. Take that out
# if you want to randomly produce a fold set on every execution
foldToMapGestureToTrial = generate_kfolds(numFolds, selected_gesture_set.map_gestures_to_trials, seed=5)
mapGestureToCorrectMatches = dict()

y_true = []
y_pred = []

gestureNamesSorted = selected_gesture_set.get_gesture_names_sorted()
for gestureName in gestureNamesSorted:
    mapGestureToCorrectMatches[gestureName] = 0

for i in range(0, len(foldToMapGestureToTrial)):
    trainingFolds = foldToMapGestureToTrial.copy()
    testFold = trainingFolds.pop(i)

    trainingData = []
    classLabels = np.array([])

    # build training data for this set of folds
    for trainingFold in trainingFolds:
        for trainingGestureName, trainingTrial in trainingFold.items():
            features = extract_features_example(trainingTrial)
            trainingData.append(features)
            classLabels = np.append(classLabels, trainingGestureName)

    # Here, we train SVM, the 'rbf' kernel is default
    # if you use rbf, need to set gamma and C parameters
    # play around with different kernels, read about them, and try them. What happens?
    # see:
    # - https://scikit-learn.org/stable/auto_examples/svm/plot_rbf_parameters.html#sphx-glr-auto-examples-svm-plot-rbf-parameter-
    # - https://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html
    clf = svm.SVC(kernel='linear') # kernel='rbf'
    clf.fit(np.array(trainingData), classLabels)

    # make predictions for this test set
    for testGestureName, testTrial in testFold.items():
        features = extract_features_example(testTrial)

        svmPrediction = clf.predict([features])
        y_true.append(testGestureName)
        y_pred.append(svmPrediction)

        if testGestureName == svmPrediction[0]:
            mapGestureToCorrectMatches[testGestureName] += 1

totalCorrectMatches = 0
print("SVM Results:\n")
for gesture in mapGestureToCorrectMatches:
    c = mapGestureToCorrectMatches[gesture]
    print("{0}: {1} ({2}%)".format(gesture, c, numFolds, c / numFolds * 100))
    totalCorrectMatches += mapGestureToCorrectMatches[gesture]

print("\nTotal SVM classifier accuracy :{0:.2f}%.format(totalCorrectMatches / numTrialsTotal * 100))

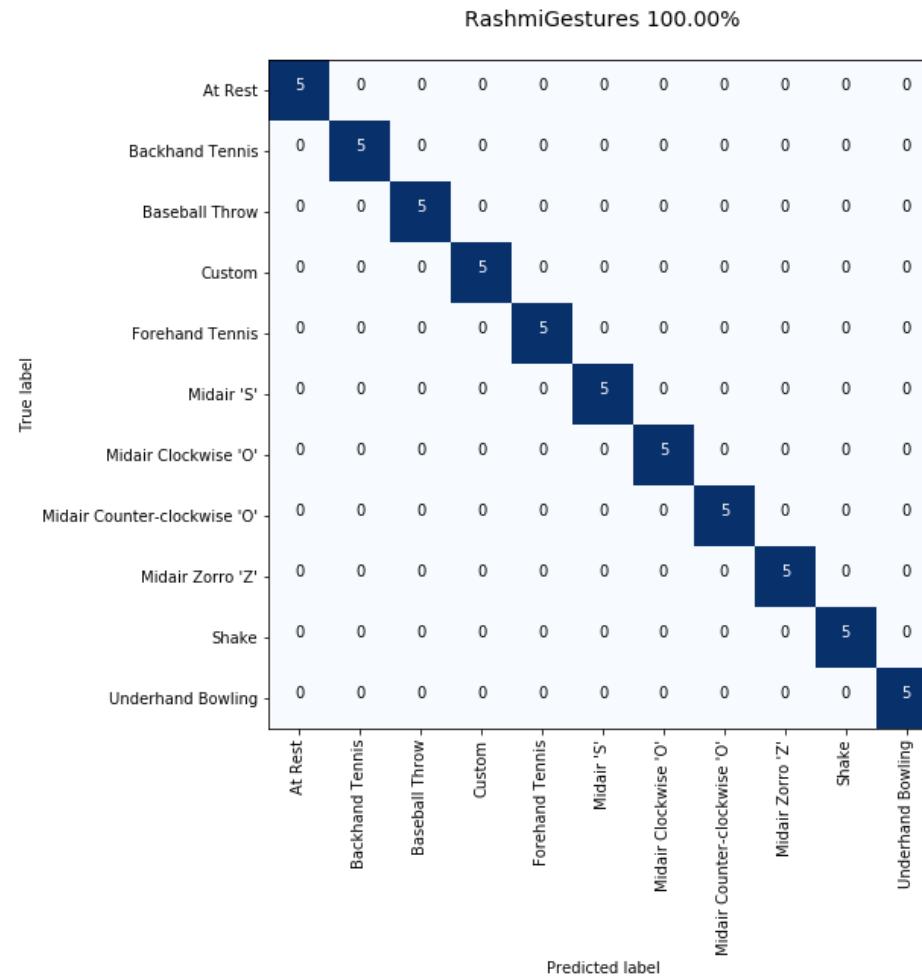
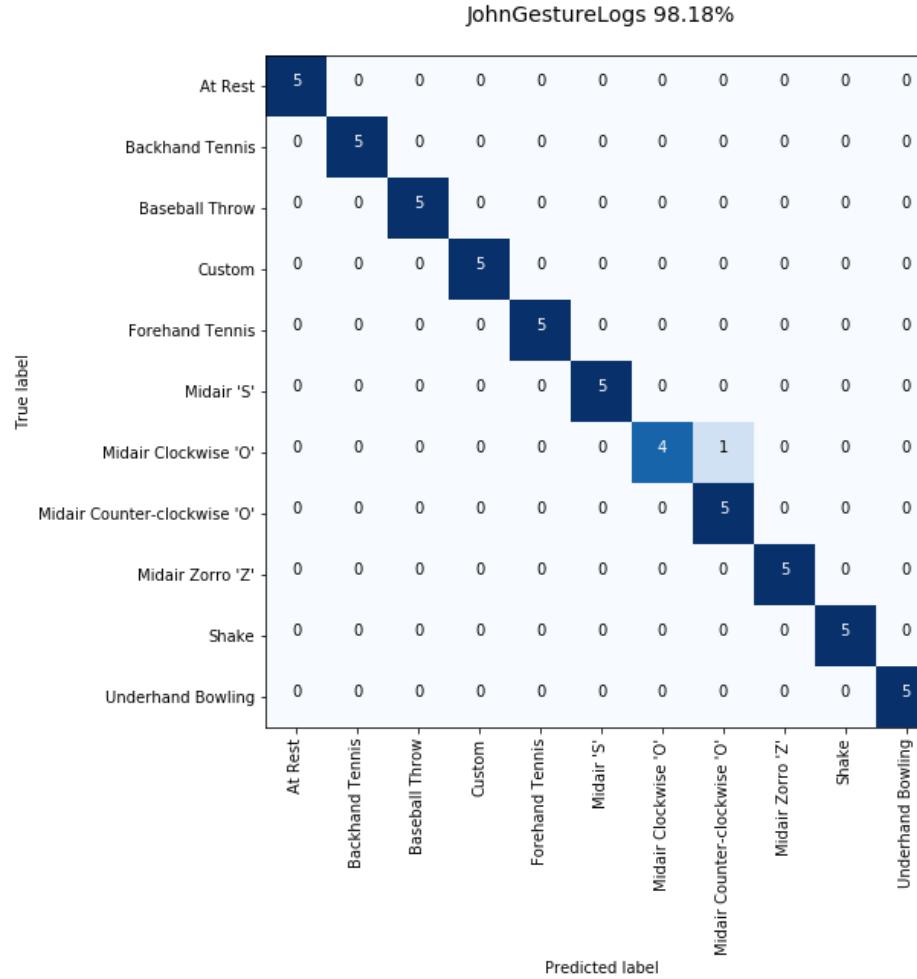
cm = confusion_matrix(y_true, y_pred, gestureNamesSorted)
plt.figure(figsize=(10,10))
plot_confusion_matrix(cm, classes=gestureNamesSorted, title='Confusion Matrix')
plt.show()
```



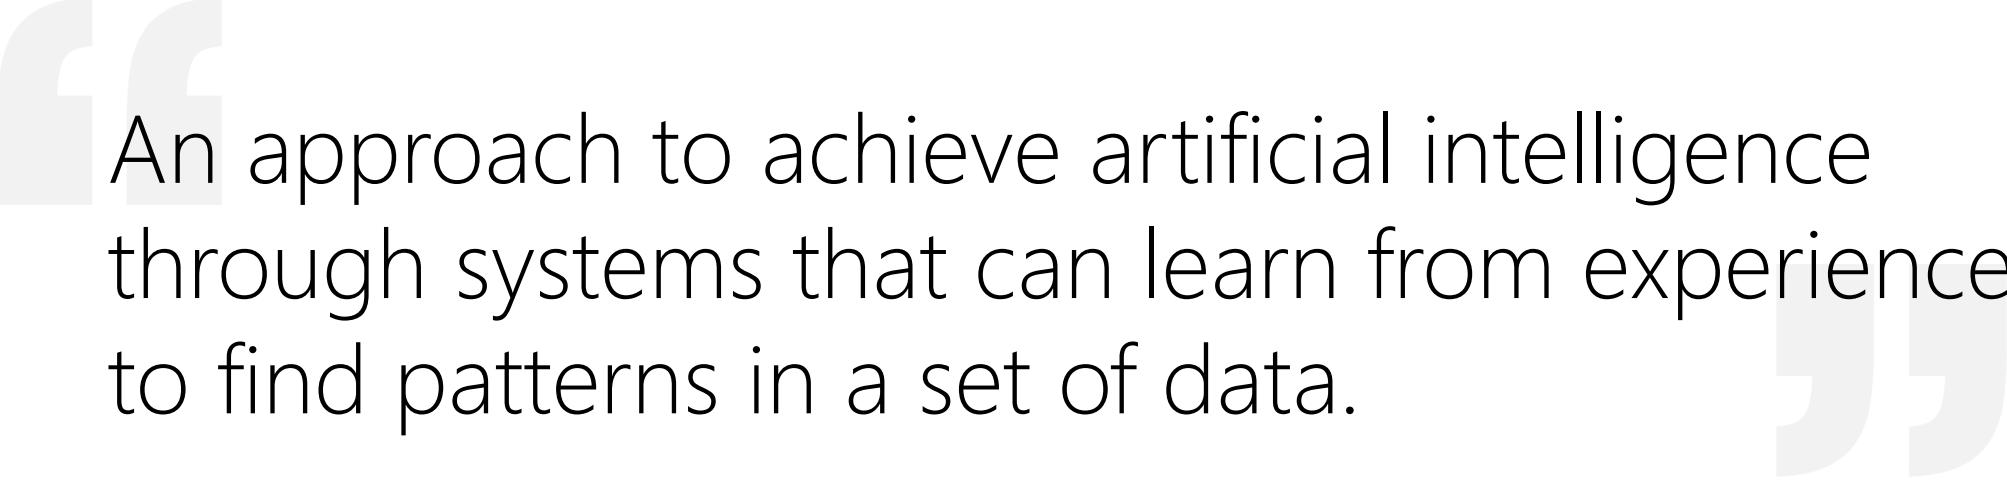
Name	Owner	Last modified	File size
AnnieGestureLogs	Anne Ross	May 15, 2019 Anne Ross	—
JohnGestureLogs	John Akers	May 15, 2019 John Akers	—
JonGestureLogs	me	May 15, 2019 me	—
JonGestureLogsHarder	me	May 15, 2019 me	—
NicoleGestures	Nicole Riley	May 15, 2019 Nicole Riley	—
RashmiGestures	Rashmi Mudduluru	May 15, 2019 Rashmi Muddulu...	—

## A4: MODEL-BASED GESTURE RECOGNIZER

# EXAMPLE SVM RESULTS ON GESTURE LOGS



# MACHINE LEARNING



An approach to achieve artificial intelligence through systems that can learn from experience to find patterns in a set of data.

**Jason Mayes**

Senior Creative Engineering, Google

# MACHINE LEARNING

The goal of machine learning is to build computer systems that can adapt and learn from their experience.

**Tom Dietterich**  
Professor, University of Oregon

# MACHINE LEARNING

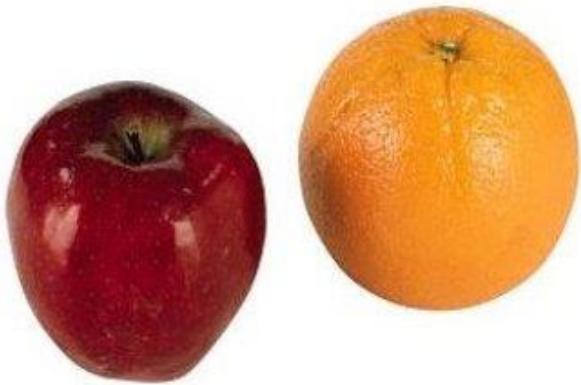
ML involves **teaching a computer to recognize patterns** by example rather than programming a system to recognize such patterns with specific rules.

In other words, **ML is about creating algorithms that learn complex models/functions from data** and then make predictions on similar data.

# SUPERVISED LEARNING

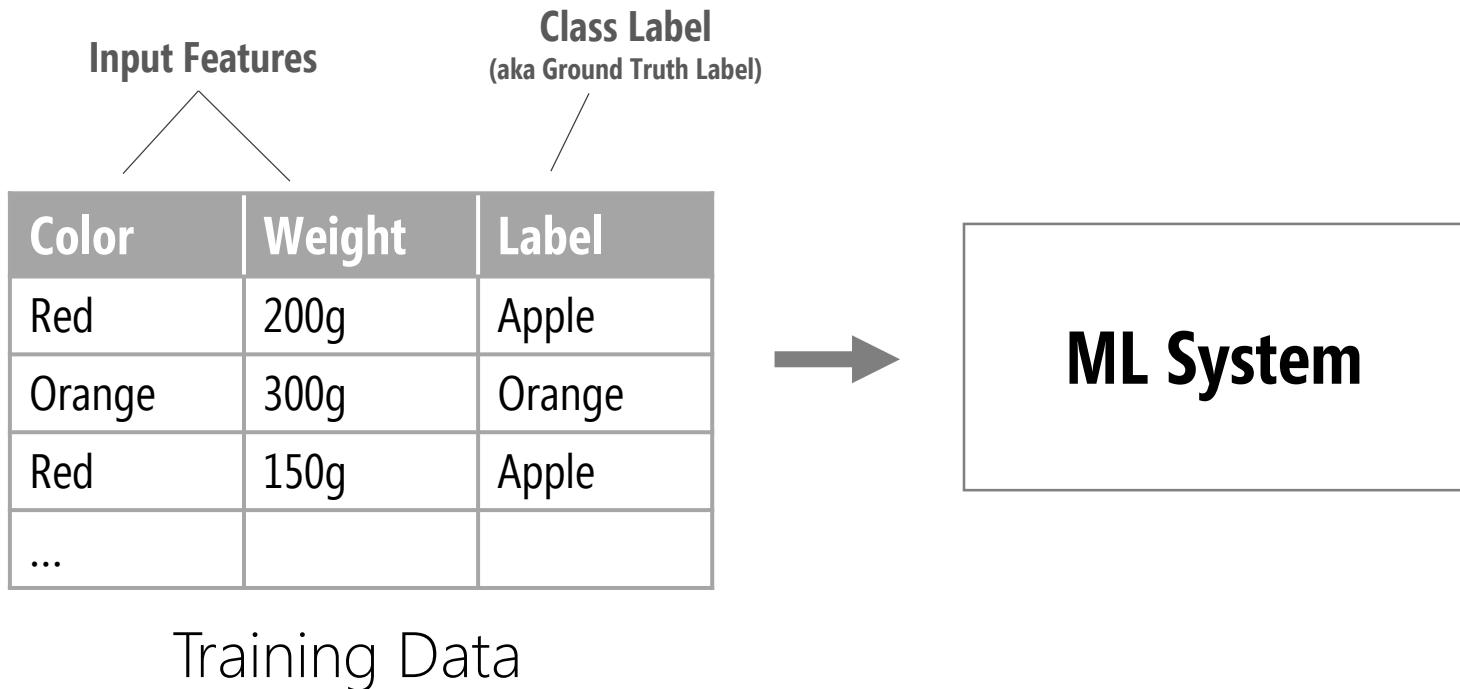
1. Input representative data to train ML system
2. ML system learns patterns from data
3. Then feed in data ML system has never seen before and it returns its best guess for what that data is

# SUPERVISED LEARNING



Example: let's build an ML model for  
recognizing apples and oranges....

# SUPERVISED LEARNING



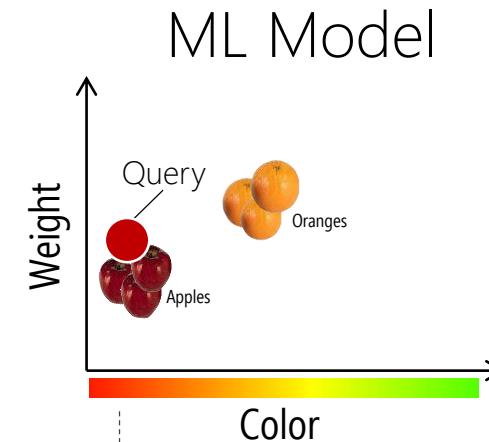
# SUPERVISED LEARNING

Input Features		Class Label (aka Ground Truth Label)
Color	Weight	Label
Red	200g	Apple
Orange	300g	Orange
Red	150g	Apple
...		

Training Data

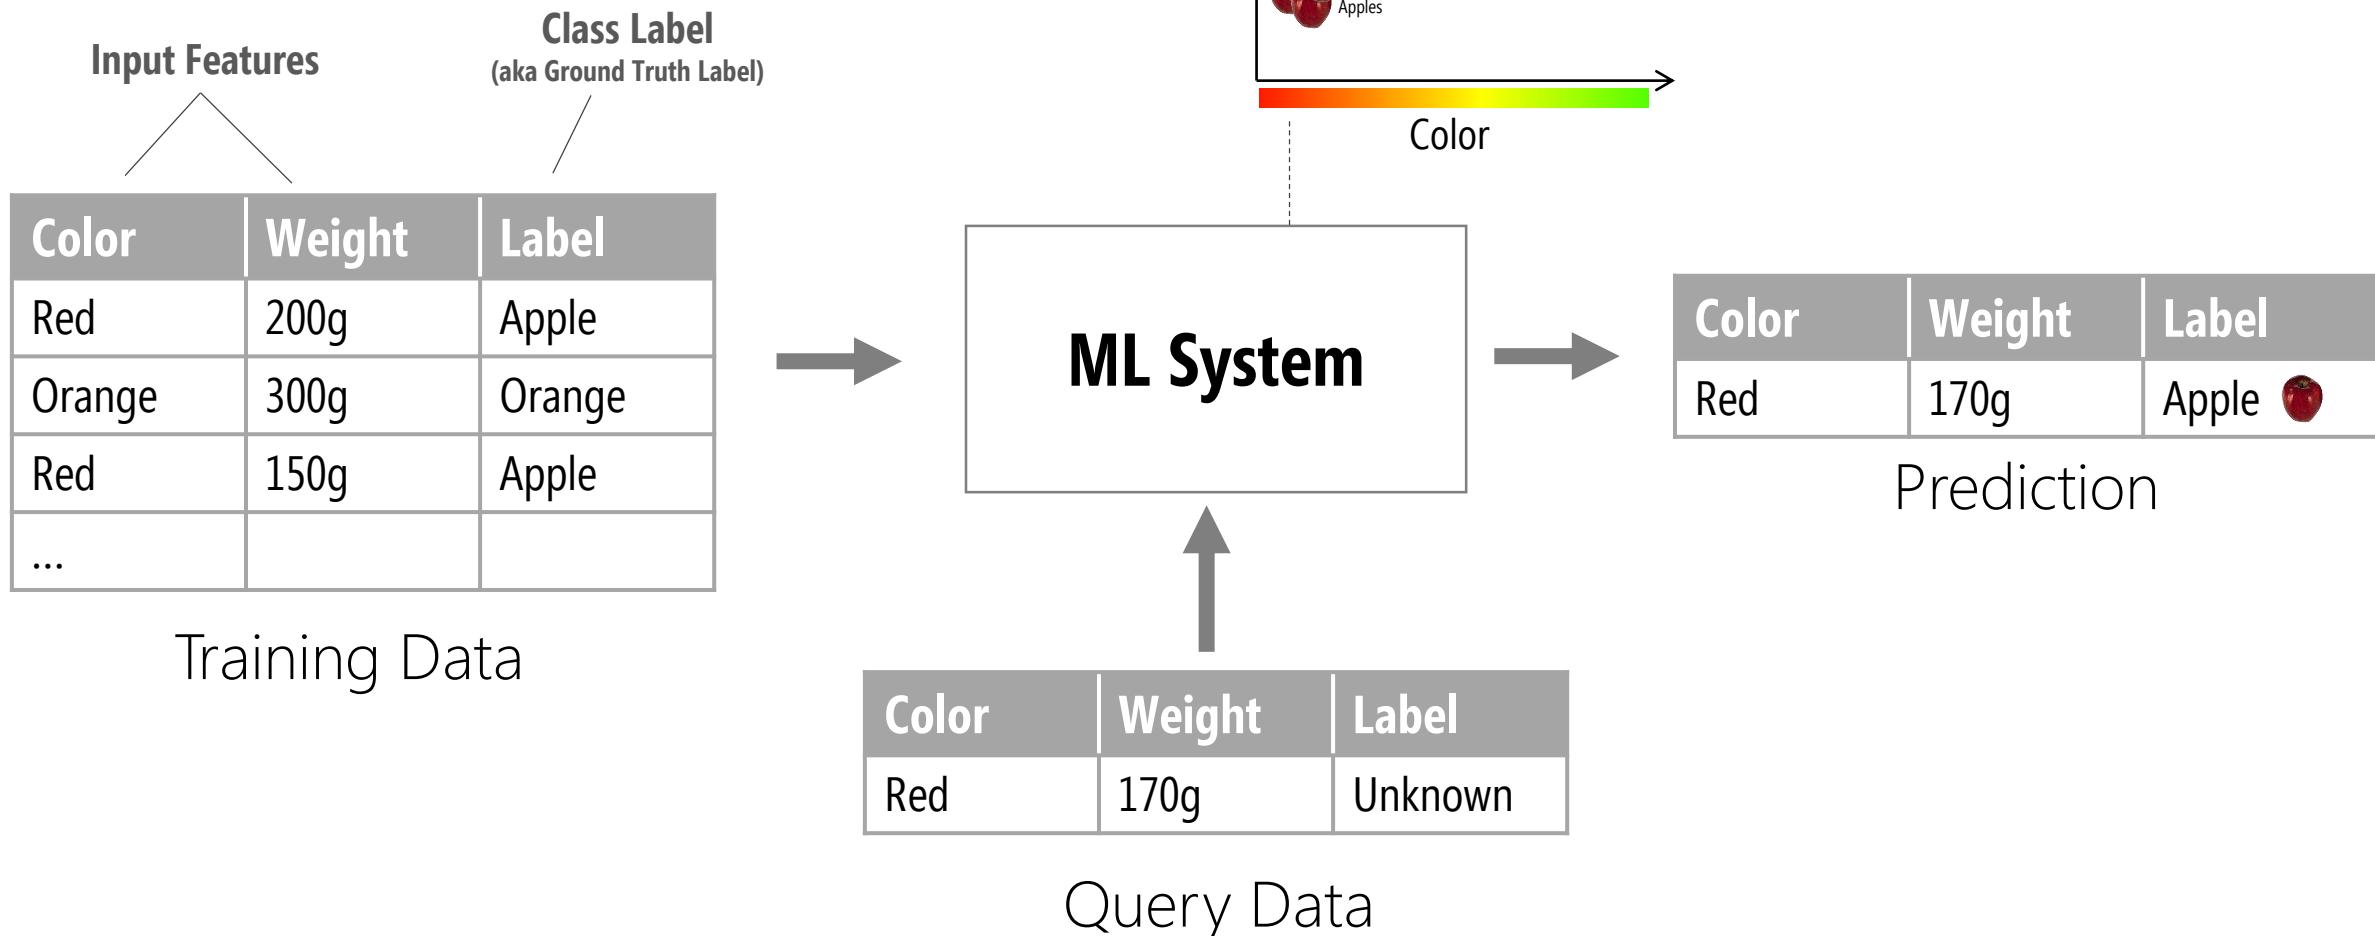
Color	Weight	Label
Red	170g	Unknown

Query Data



ML System

# SUPERVISED LEARNING



# SUPERVISED LEARNING PROBLEMS AND APPROACHES

## CLASSIFICATION PROBLEMS

When you need to classify some input vector into a category (*e.g.*, Apple or Orange). Output is some class label  $C$ .

### Popular approaches:

Support Vector Machines (SVMs)

Random forests

Neural Networks

Deep Learning

## REGRESSION PROBLEMS

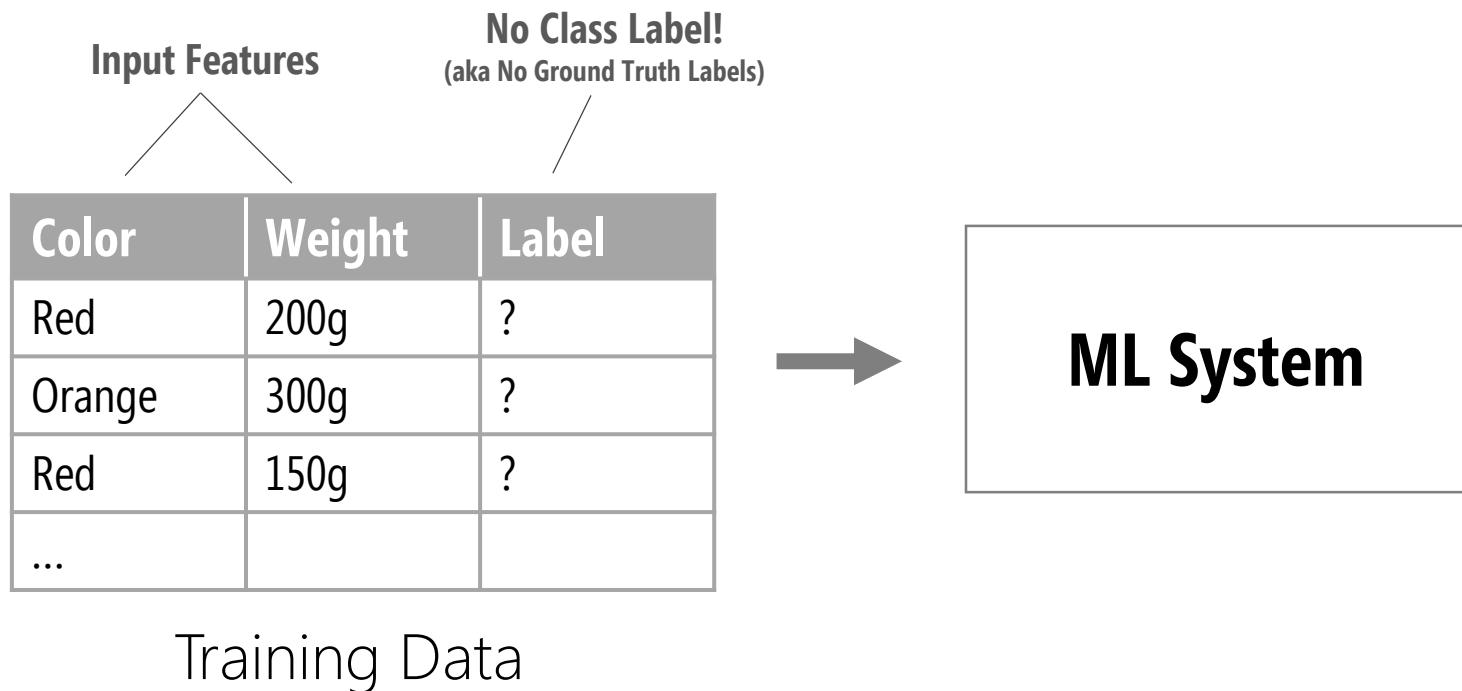
When you want to predict some numeric value (*e.g.*, house prices next year in Seattle). Output has infinitely many values (a number).

### Popular approaches:

Linear regression

Random forests

# UNSUPERVISED LEARNING



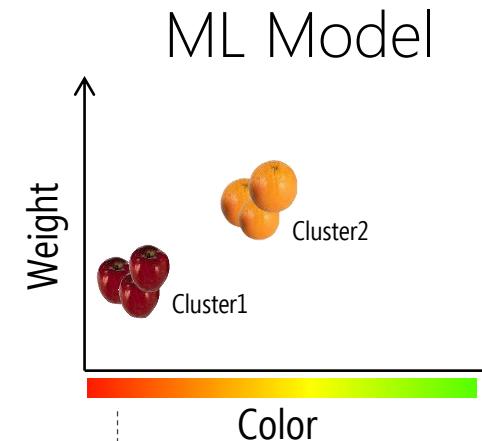
# UNSUPERVISED LEARNING

**Input Features**

**No Class Label!**  
(aka No Ground Truth Labels)

Color	Weight	Label
Red	200g	?
Orange	300g	?
Red	150g	?
...		

Training Data



# UNSUPERVISED LEARNING

**Input Features**

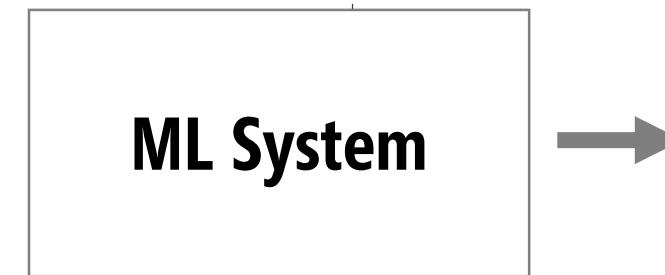
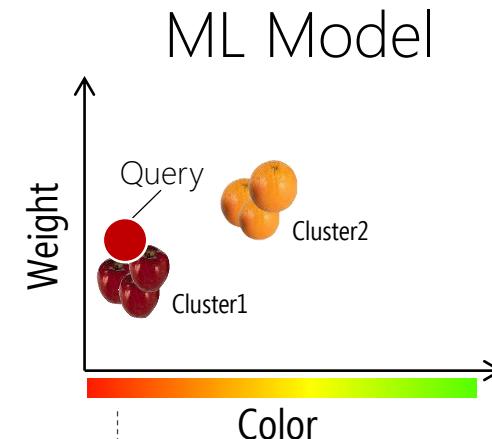
**No Class Label!**  
(aka No Ground Truth Labels)

Color	Weight	Label
Red	200g	?
Orange	300g	?
Red	150g	?
...		

Training Data

Color	Weight	Label
Red	170g	?

Query Data



Color	Weight	Label
Red	170g	Cluster1

Prediction

# UNSUPERVISED LEARNING PROBLEMS AND APPROACHES

## CLUSTERING PROBLEMS

When you want to discover inherent groupings in your data (*e.g.*, grouping customers by purchasing behavior)

### Popular approaches:

K-means clustering

## ASSOCIATION PROBLEMS

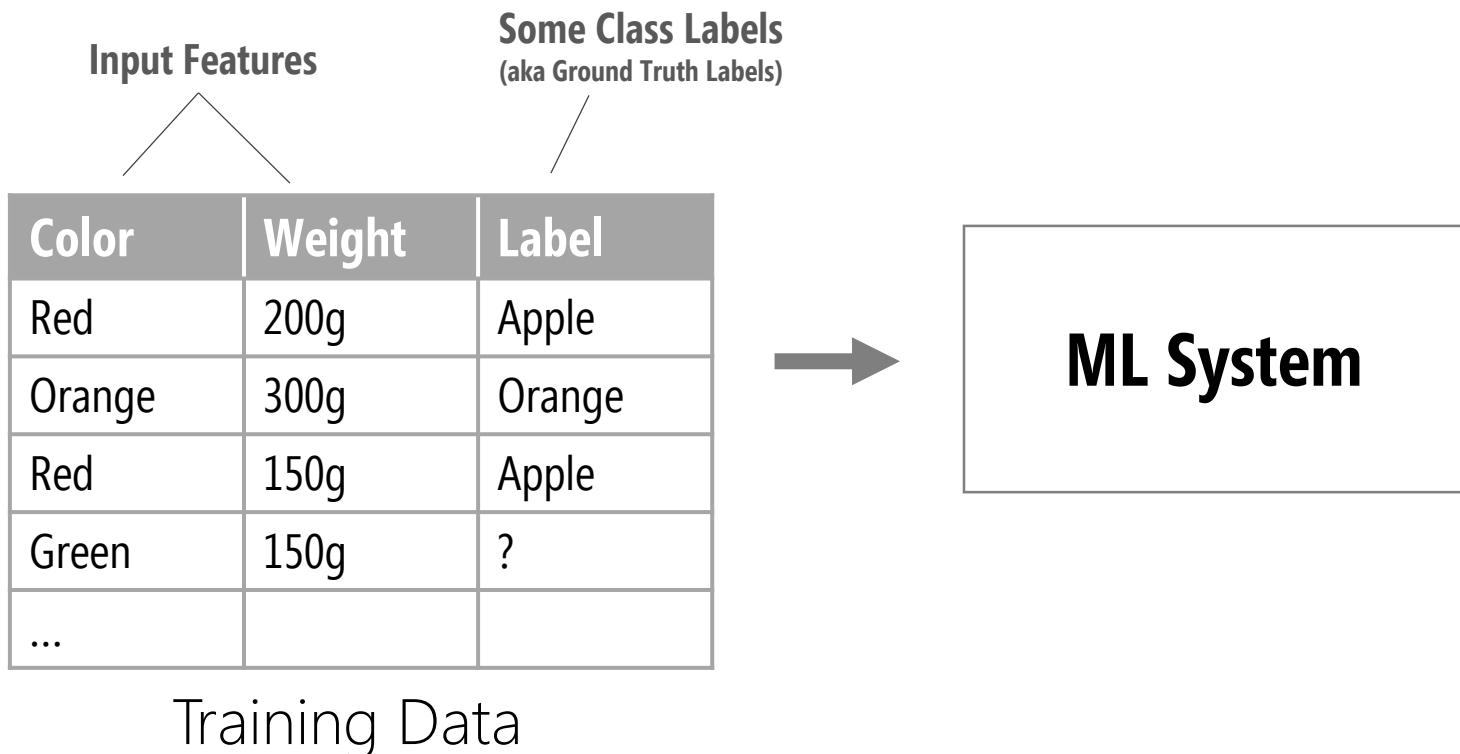
When you want to discover associative rules between large portions of data (*e.g.*, people that buy X also tend to buy Y)

### Popular approaches:

Apriori algorithm

Collaborative filtering

# SEMISUPERVISED LEARNING



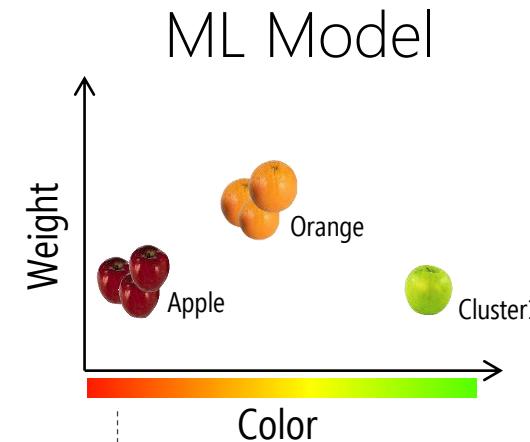
# SEMISUPERVISED LEARNING

Input Features

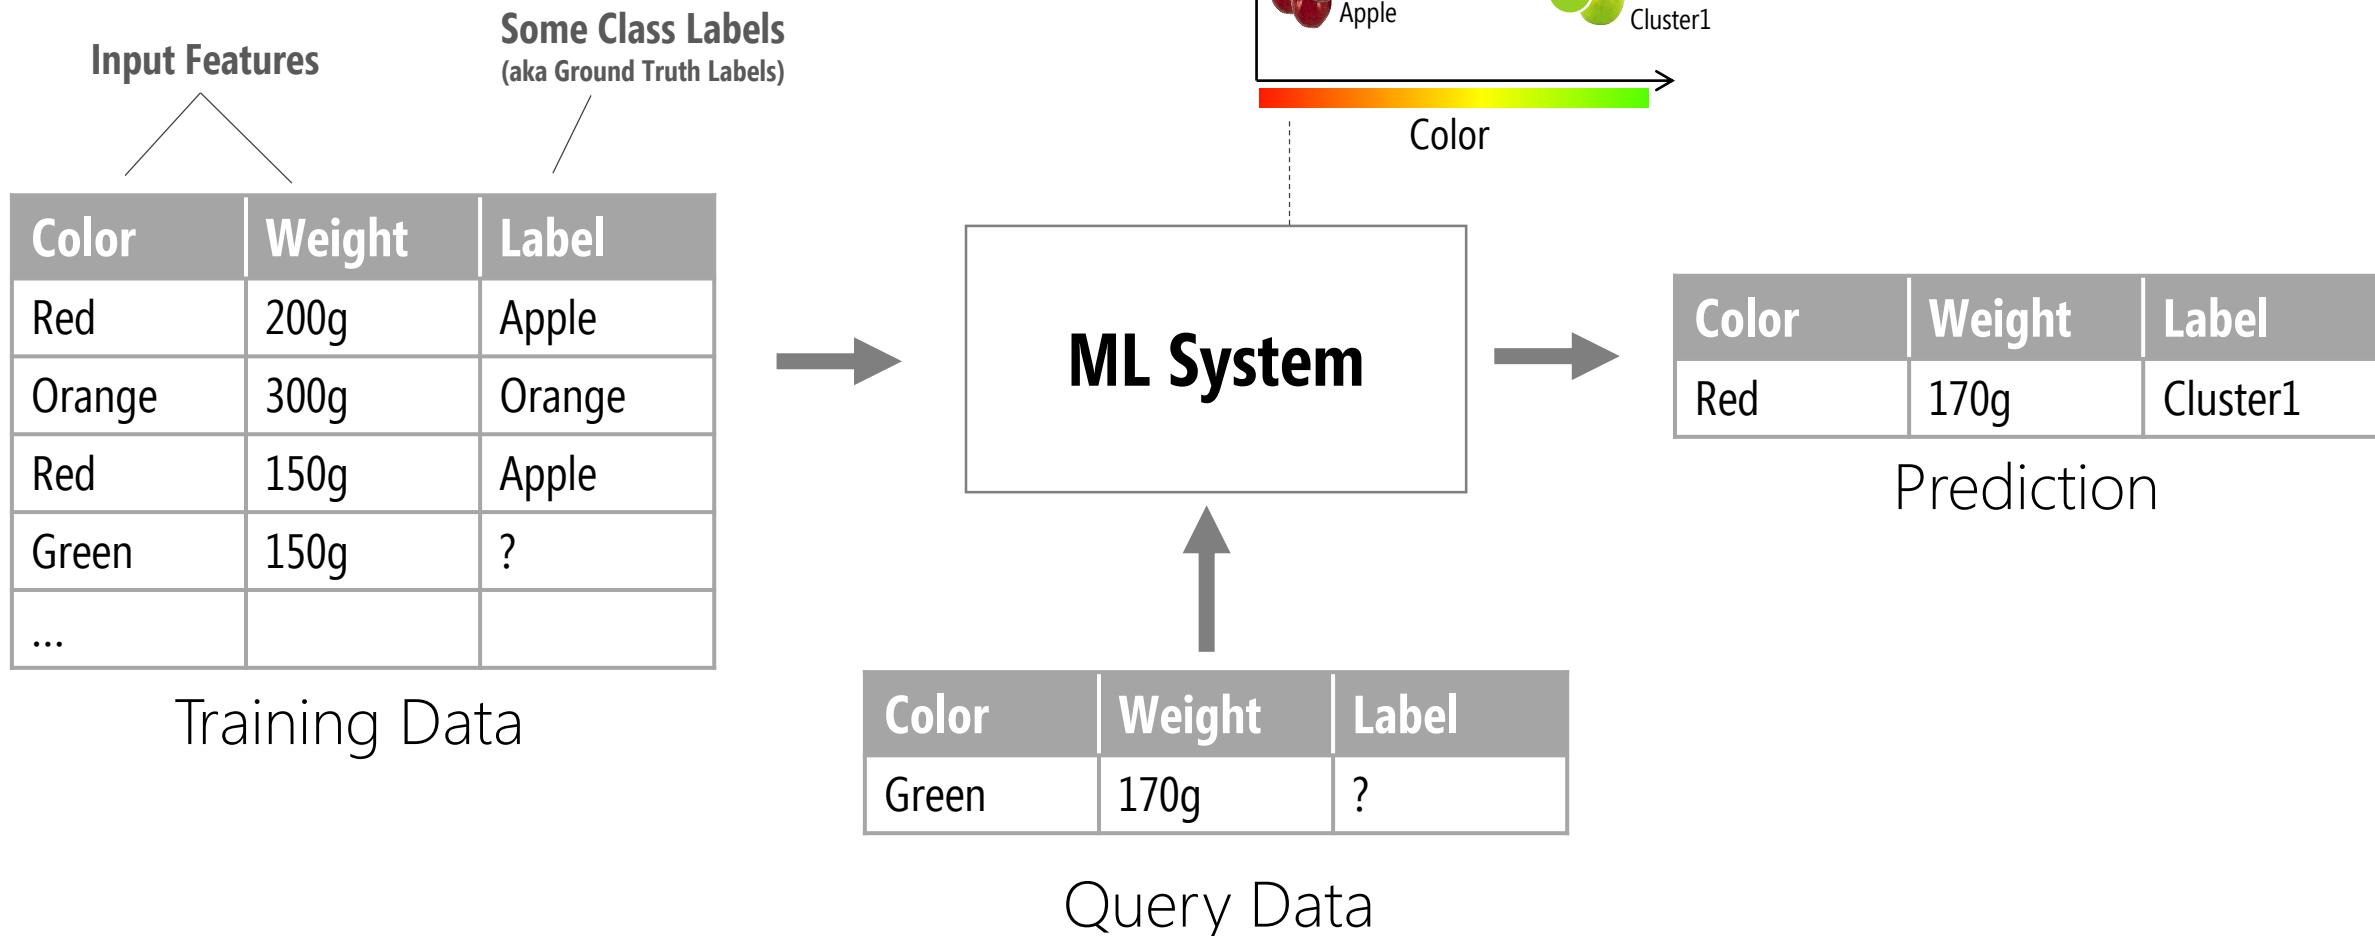
Some Class Labels  
(aka Ground Truth Labels)

Color	Weight	Label
Red	200g	Apple
Orange	300g	Orange
Red	150g	Apple
Green	150g	?
...		

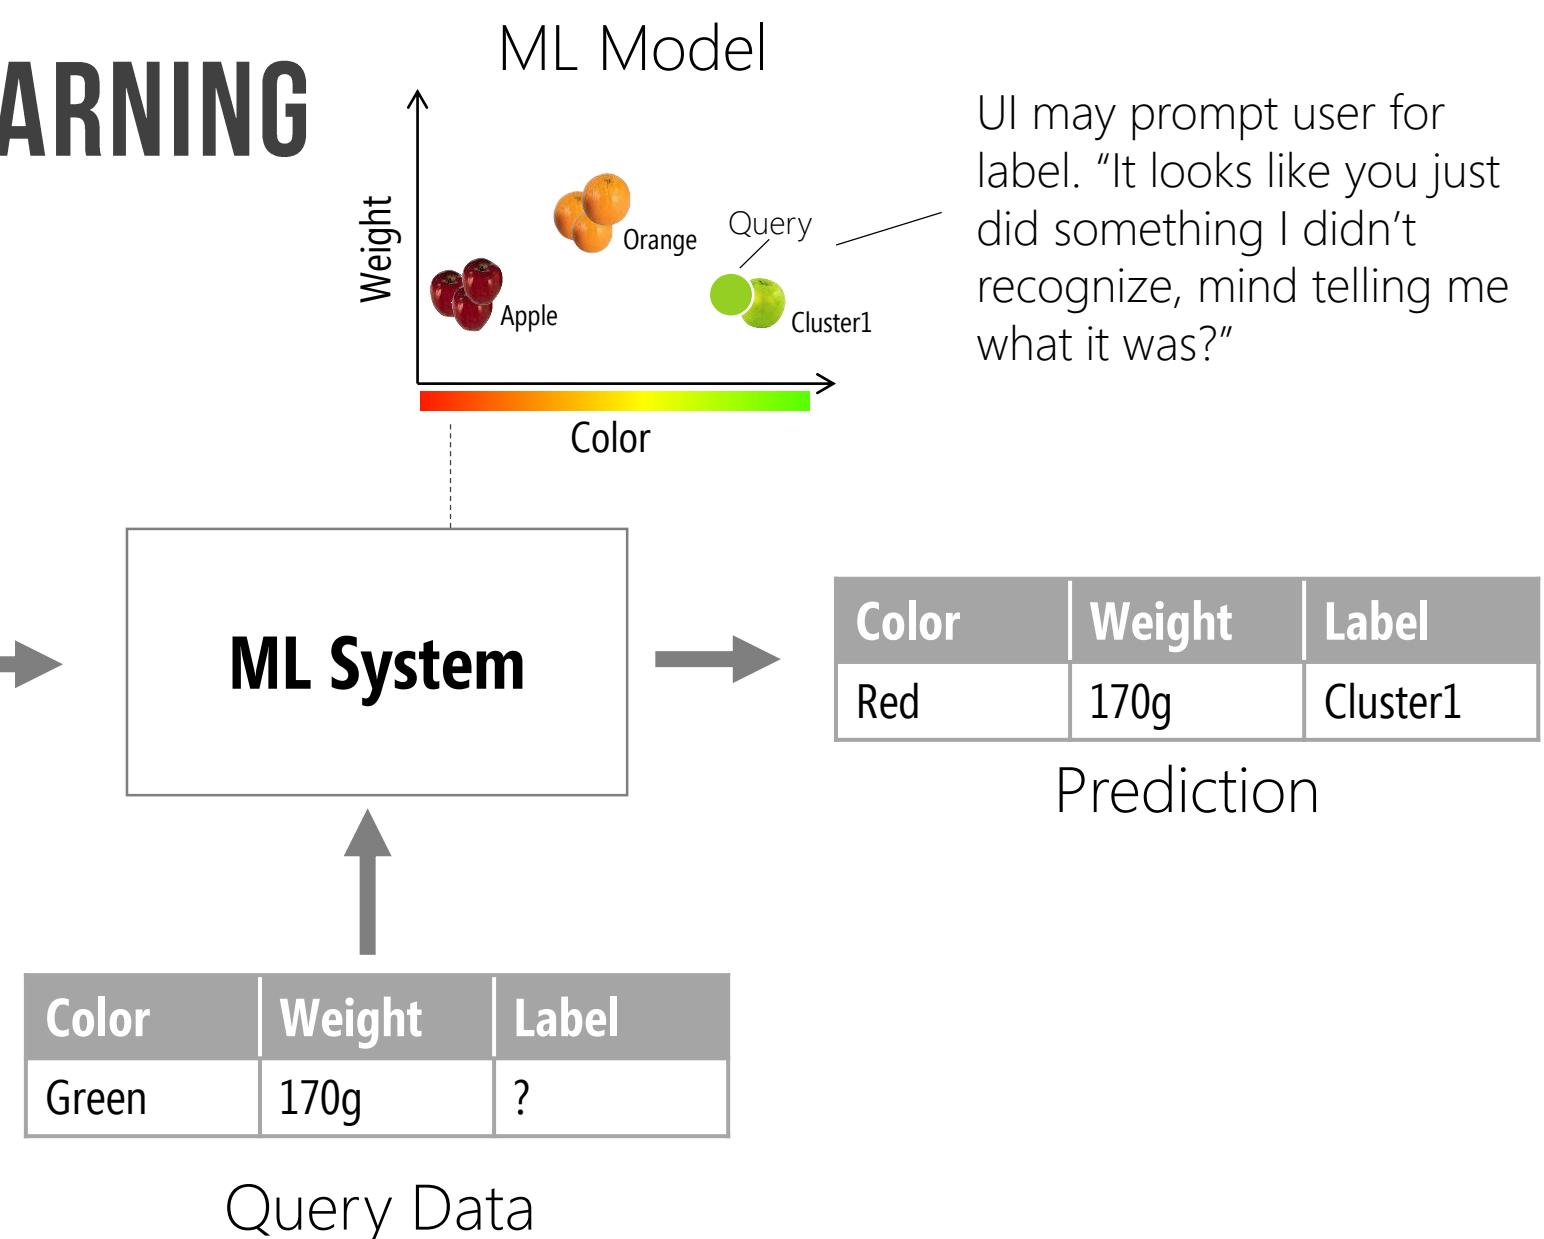
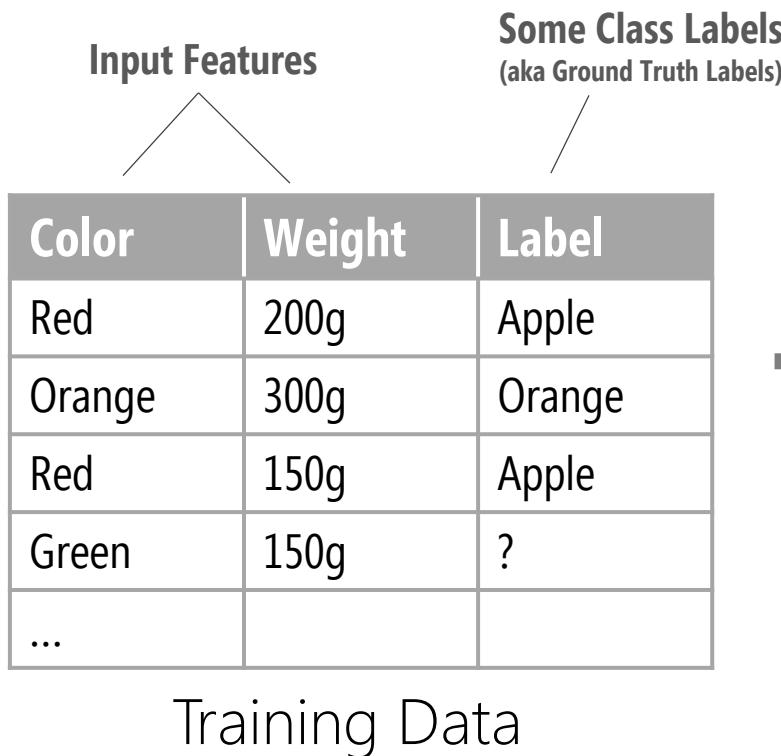
Training Data



# SEMISUPERVISED LEARNING



# SEMISUPERVISED LEARNING



# HOW TO CHOOSE A MODEL?

Your decision should be based on the nature of the dataset (e.g., images, videos, text, continuous time-series)

And the question you're trying to answer (e.g., What is likely to happen next? What are the anomalies in this dataset? What objects are in this image? What gesture is the user performing?)

Often, to answer these questions, we explore the literature to see how others have solved similar problems and adapt...

For A2, we strongly recommend an SVM

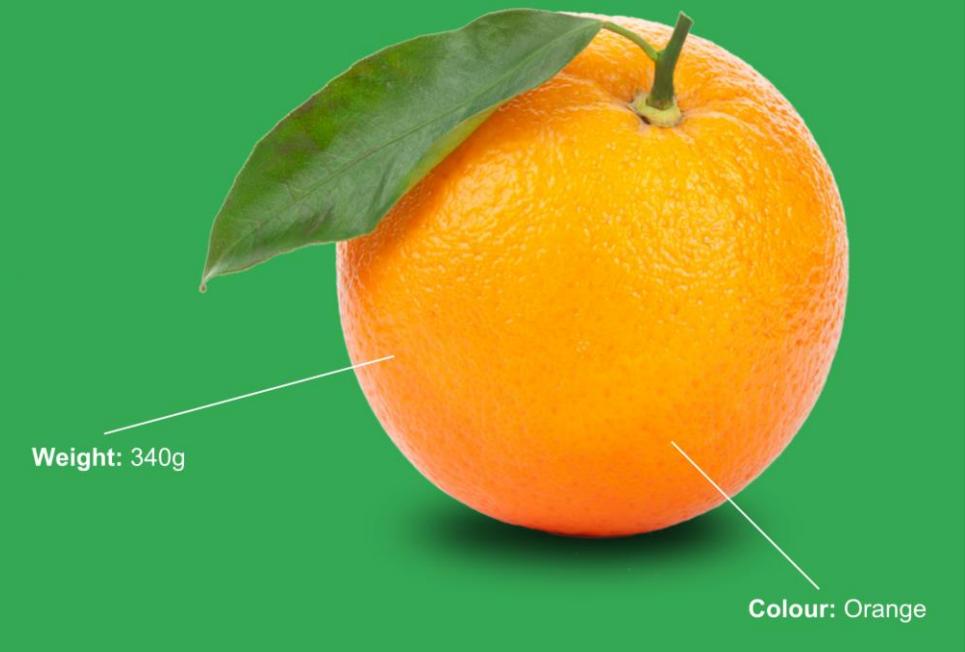
**How** to select our **input features?**

# FEATURES

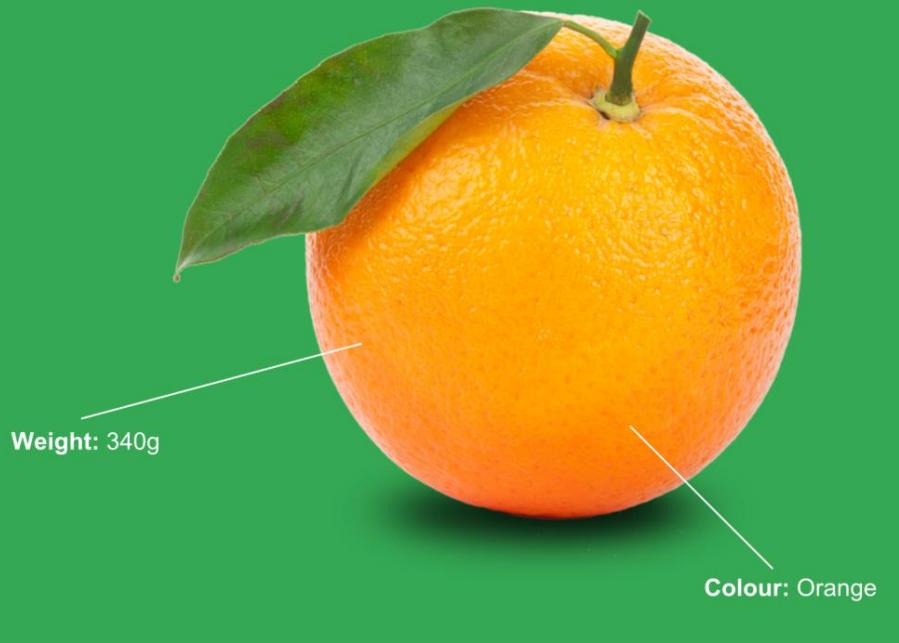
Features are used to **train** the ML system.

You want to **pick features that you think help discriminate** your dataset.

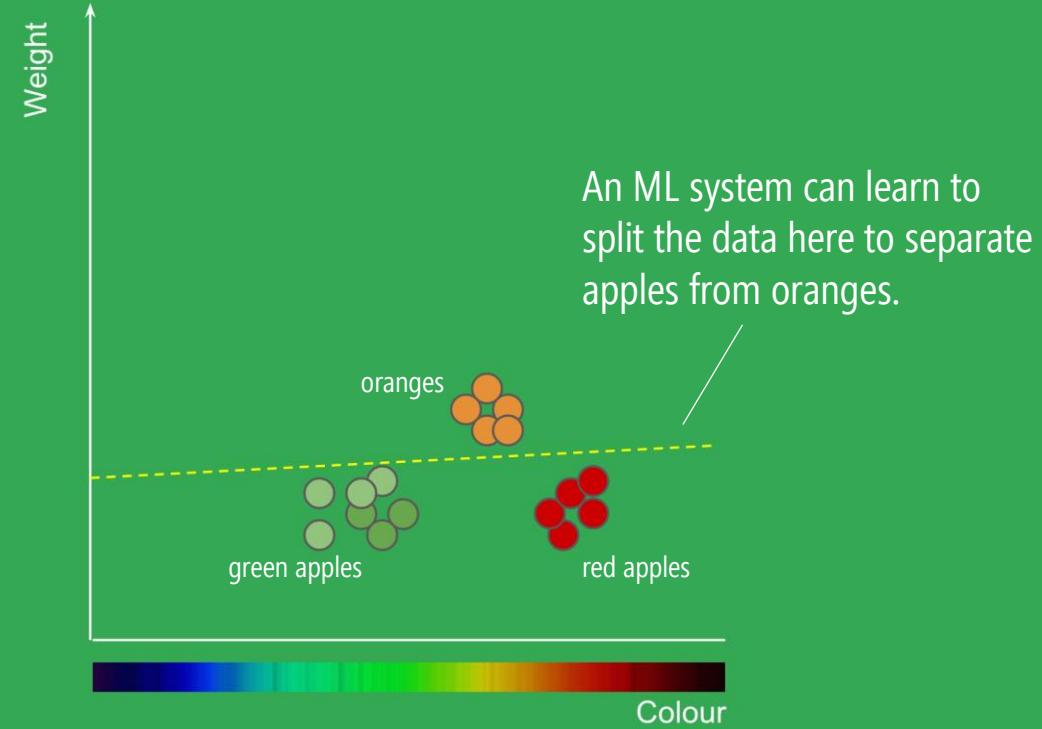
With the fruit example, we tried just **weight** and **color**. **Two features**, which means your featureset is **two-dimensional**



# SELECTING FEATURES

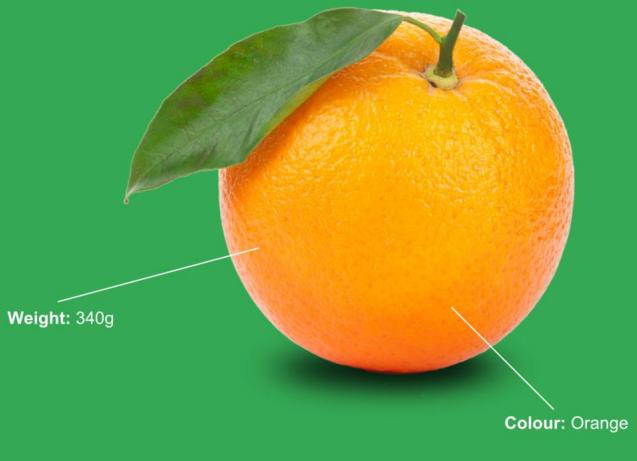


**Picking features.** When we were building the classifier to automatically distinguish fruits, we setup a 2D feature vector: weight and color



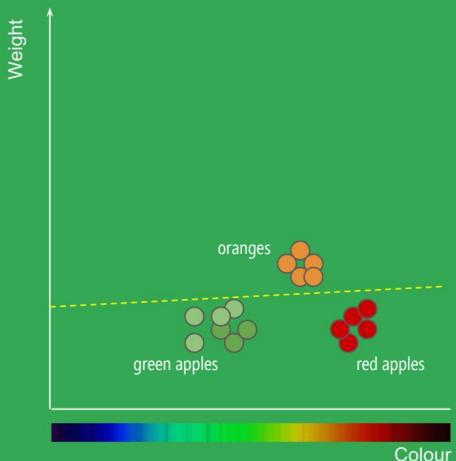
**Brainstorm, visualize, investigate, & iterate your feature space.** When exploring features, it's a good idea to plot them on a graph (when possible) to develop intuition & assess discriminability.

# SELECTING FEATURES

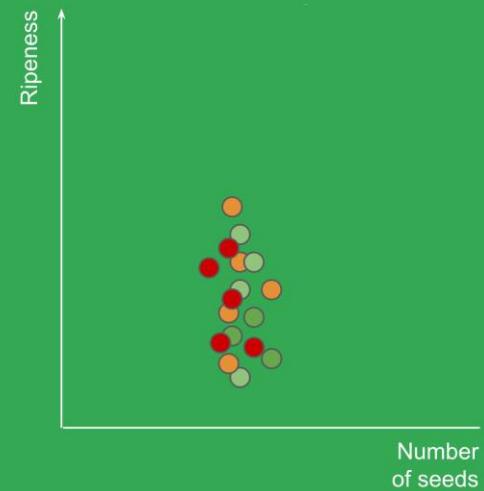


**Picking features.** if we were building a classifier to automatically distinguish fruits, we could setup a 2D feature vector: weight and color

**Brainstorm, visualize, investigate, & iterate your feature space.** When exploring features, it's a good idea to plot them on a graph (when possible) to develop intuition and understand their discriminability.



Ripeness and # of seeds are poor features for distinguishing oranges and apples



**Choose your features carefully.** Some features may not help you distinguish data. For example, ripeness and number of seeds are bad features.

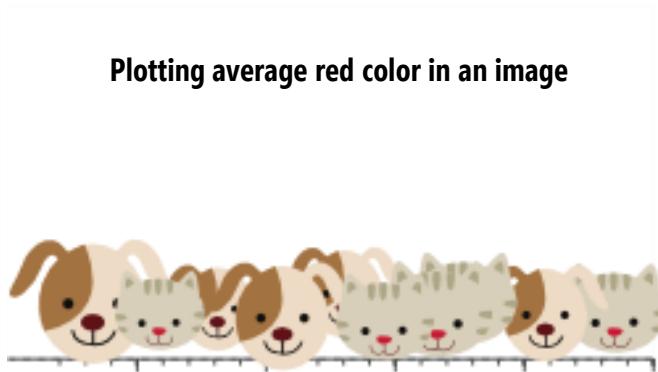
# SELECTING FEATURES: HIGHER DIMENSIONALITY

Generally, ML systems are trained with (super!) high dimensional feature vectors that are hard (impossible) to plot and develop an intuition for.

We can plot maybe ~4 dimensions at a time (x, y, z-axis + some visual attribute of the marker).

# SELECTING FEATURES: HIGHER DIMENSIONALITY

Let's look at another example: this time trying to separate cats and dogs using average pixel color in image.



## 1D Feature Vector

Plotting average red color in image.  
Notice how a single feature does not result in a perfect separation of our training data.

# EXERCISE: BRAINSTORM & PLOT FEATURES

We are going to be doing exercises in Jupyter Notebook again.

The screenshot shows a GitHub repository page for [jonfroehlich/CSE599Sp2019/tree/master/Assignments/A4-SVMGestureRecognizer](https://github.com/jonfroehlich/CSE599Sp2019/tree/master/Assignments/A4-SVMGestureRecognizer). The repository has 7 stars, 5 forks, and 0 open issues. The commit history shows a single commit from [jonfroehlich](#) adding a basic SVM gesture recognizer for the assignment, made 3 hours ago. The commit message is "Added basic SVM gesture recognizer for assignment".

GitHub, Inc. [US] | https://github.com/jonfroehlich/CSE599Sp2019/tree/master/Assignments/A4-SVMGestureRecognizer

Pull requests Issues Marketplace Explore

jonfroehlich / CSE599Sp2019

Unwatch 7 Star 5 Fork 0

Code Issues 1 Pull requests 0 Projects 0 Wiki Insights Settings

Branch: master CSE599Sp2019 / Assignments / A4-SVMGestureRecognizer / Create new file Upload files Find file History

jonfroehlich Added basic SVM gesture recognizer for assignment Latest commit a4ab214 3 hours ago

SVMGestureRecognizer.ipynb Added basic SVM gesture recognizer for assignment 3 hours ago

SVMGestureRecognizer x +

localhost:8889/notebooks/SVMGestureRecognizer.ipynb

jupyter SVMGestureRecognizer (unsaved changes)

File Edit View Insert Cell Kernel Widgets Help

Not Trusted Python 3

Table of Contents

- Major Classes and File Parsing
- Load the Data
- Signal Exploration
- SVM Experiment Infrastructure
- The SVM

```
In [1]: # This cell includes the major classes used in our classification analyses
import matplotlib.pyplot as plt # needed for plotting
import numpy as np # numpy is primary library for numeric array (and matrix) handling
import scipy as sp
from scipy import stats, signal
import random
from sklearn import svm # needed for svm
from sklearn.metrics import confusion_matrix
import itertools

# Each accelerometer log file gets parsed and made into a SensorData object
class SensorData:

    # Constructors in Python Look Like this (strangely enough)
    # ALL arguments are numpy arrays except sensorType, which is a str
    def __init__(self, sensorType, currentTimeMs, sensorTimestampMs, x, y, z):
        self.sensorType = sensorType

    # On my mac, I could cast as straight-up int but on Windows, this failed
    # This is because on Windows, a Long is 32 bit but on Unix, a Long is 64bit
    # So, forcing to int64 to be safe. See: https://stackoverflow.com/q/38314118
    self.currentTimeMs = currentTimeMs.astype(np.int64)

    # sensorTimestampMs comes from the Arduino function
    # https://www.arduino.cc/en/Language/functions/time/millis/
    # which returns the number of milliseconds passed since the Arduino board began running the current program.
    self.sensorTimestampMs = sensorTimestampMs.astype(np.int64)

    self.x = x.astype(float)
    self.y = y.astype(float)
    self.z = z.astype(float)

    # Calculate the magnitude of the signal
    self.mag = np.sqrt(self.x**2 + self.y**2 + self.z**2)

    self.sampleLengthInSecs = (self.currentTimeMs[-1] - self.currentTimeMs[0]) / 1000.0
    self.samplesPerSecond = len(self.currentTimeMs) / self.sampleLengthInSecs

    # Returns a dict of numpy arrays for each axis of the accel + magnitude
    def get_data(self):
        return {"x":self.x, "y":self.y, "z":self.z, "mag":self.mag}

    # Returns a dict of numpy arrays for each axis of the accel + magnitude
    def get_processed_data(self):
        return {"x_p":self.x_p, "y_p":self.y_p, "z_p":self.z_p, "mag_p":self.mag_p}

    # A trial is one gesture recording and includes an accel SensorData object
    # In the future, this could be expanded to include other recorded sensors (e.g., a gyro)
    # that may be recorded simultaneously
    class Trial:

        # We actually parse the sensor log files in the constructor--this is probably bad practice
        # But offers a relatively clean solution
        def __init__(self, gestureName, endTimeMs, trialNum, accelLogFilenameWithPath):
            self.gestureName = gestureName
```

# EXERCISE: BRAINSTORM & PLOT FEATURES

What's the best 1D feature vector you can think of for recognizing gestures?

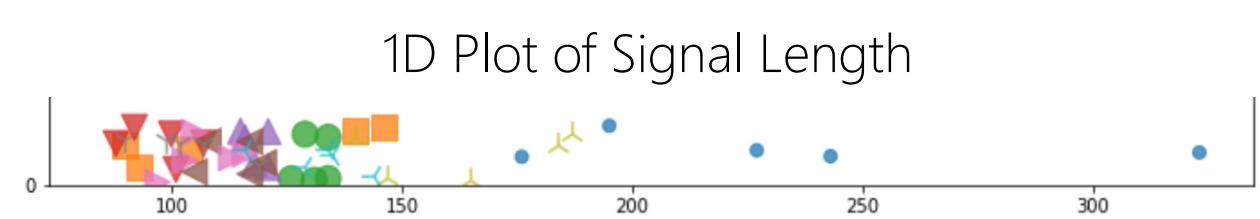
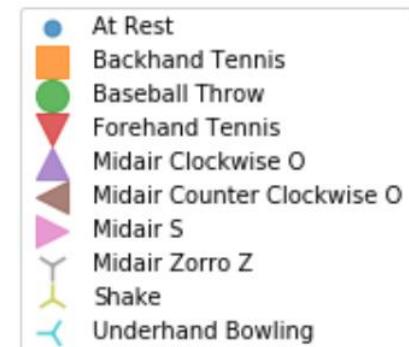
Length of signal

Max accel magnitude

Fundamental frequency of FFT

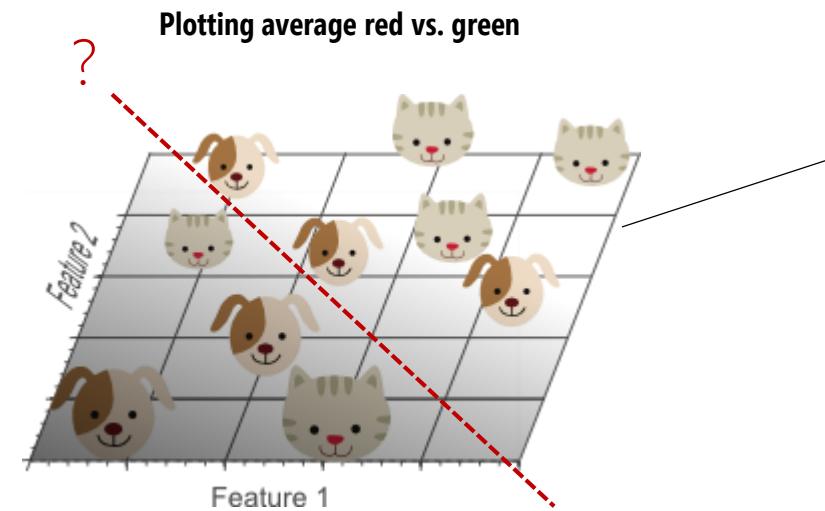
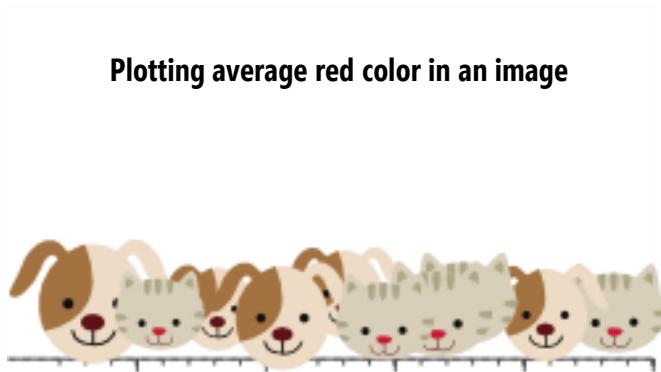
...

Plot your ideas on individual scatterplot graphs; color the gesture classes uniquely (you'll make a 2D scatter graph where the y-values are just small random perturbations to view your data better). We want to see **classes grouped together** and a **clear separation between them**.



# SELECTING FEATURES: HIGHER DIMENSIONALITY

Let's look at another example: this time trying to separate cats and dogs using average pixel color in image.



## 1D Feature Vector

Plotting average red color in image.  
Notice how a single feature does not result in a perfect separation of our training data.

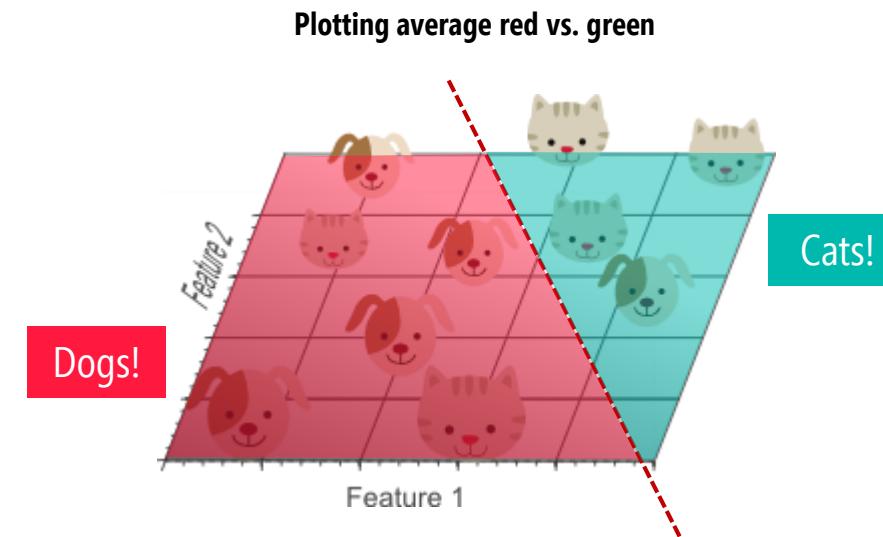
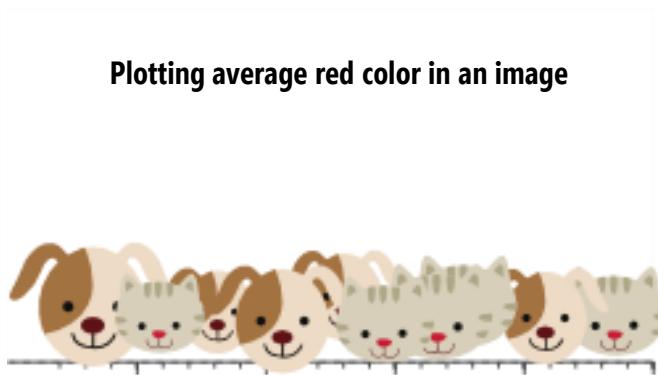
## 2D Feature Vector

Plotting average red and green colors in image. Now we begin to see some more separation but not perfect.

Is there any straight line that we can draw through our data that separates cats from dogs?

# SELECTING FEATURES: HIGHER DIMENSIONALITY

Let's look at another example: this time trying to separate cats and dogs using average pixel color in image.



**Overall Accuracy: 7/10 (70%)**

True	Cat	Predicted		Overall Accuracy
		Cat	Dog	
Cat	3	2		3/5 (60%)
Dog	1	4		4/5 (80%)

## 1D Feature Vector

Plotting average red color in image.  
Notice how a single feature does not result in a perfect separation of our training data.

## 2D Feature Vector

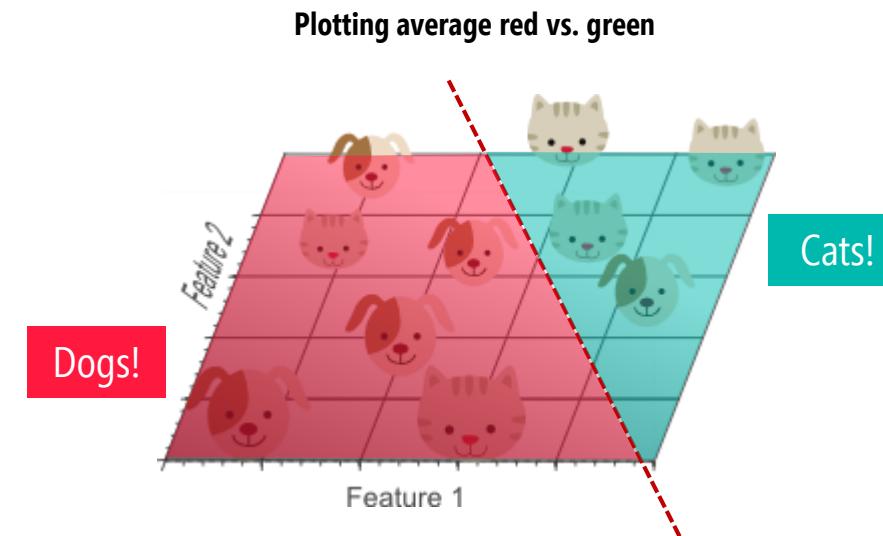
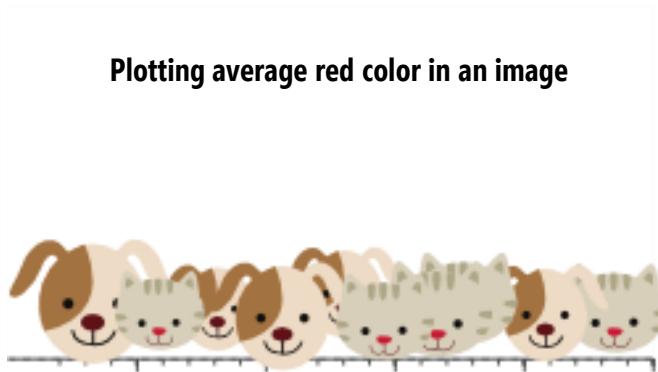
Plotting average red and green colors in image. Now we begin to see some more separation but not perfect.

## Results

Overall accuracy is 70% with more cats being misclassified as dogs.

# SELECTING FEATURES: HIGHER DIMENSIONALITY

Let's look at another example: this time trying to separate cats and dogs using average pixel color in image.



## 1D Feature Vector

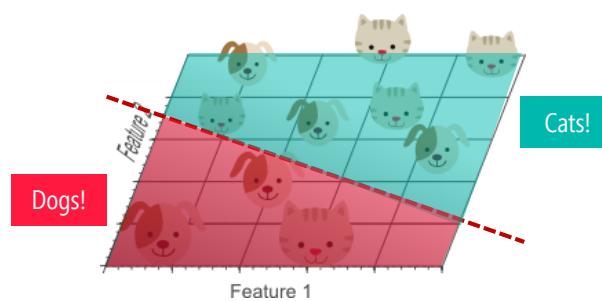
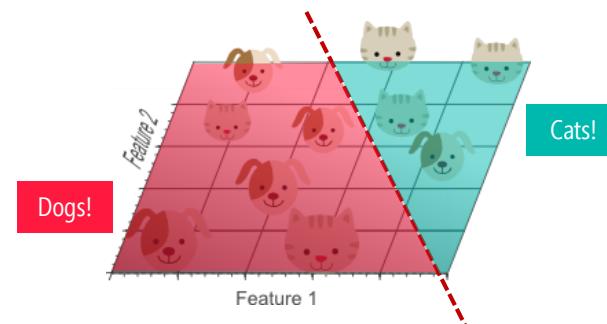
Plotting average red color in image.  
Notice how a single feature does not result in a perfect separation of our training data.

## 2D Feature Vector

Plotting average red and green colors in image. Now we begin to see some more separation but not perfect.

# SELECTING FEATURES: HIGHER DIMENSIONALITY

And, of course, you can find different splits of the data. But we want to minimize error—an optimization problem—to find best split.



**Overall Accuracy: 7/10 (70%)**

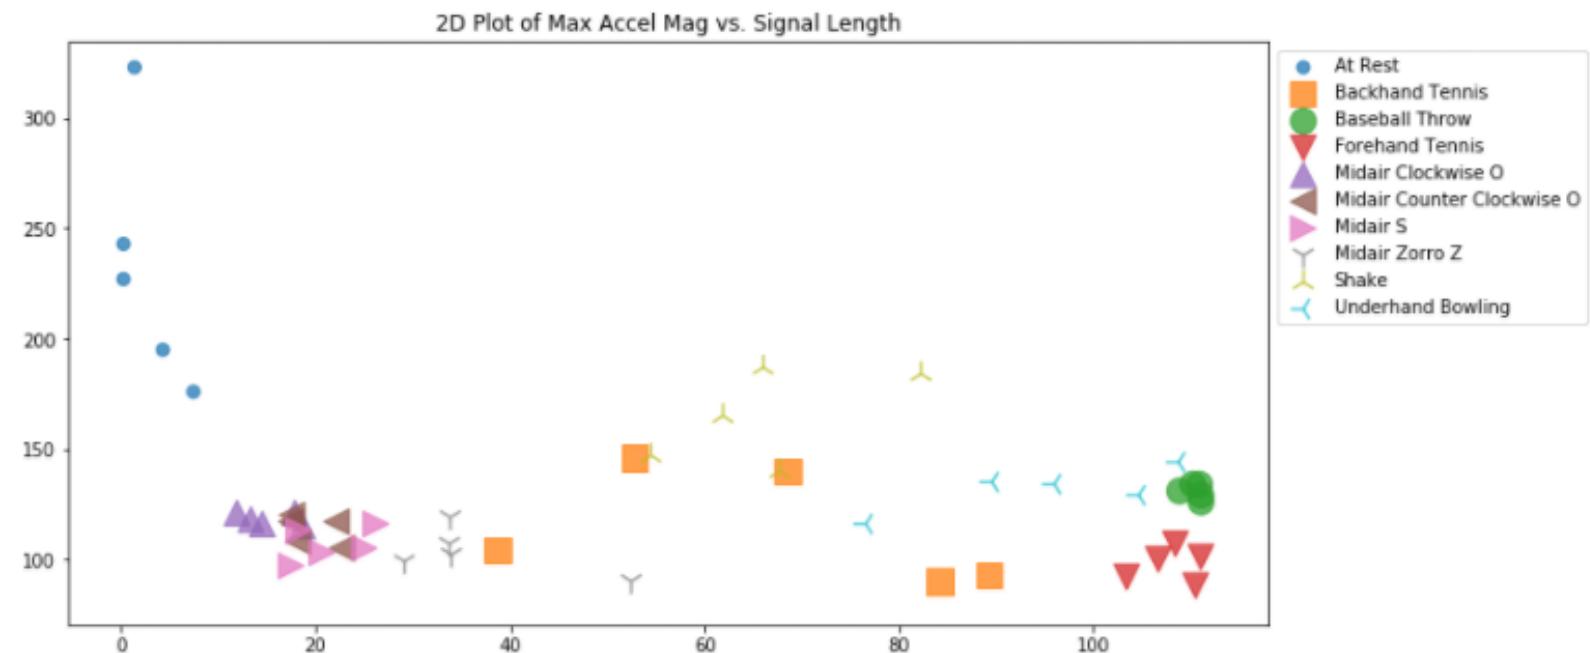
True			Overall Accuracy: 7/10 (70%)
	Cat	Dog	
Cat	3	2	3/5 (60%)
	1	4	4/5 (80%)
Cat		Dog	
Predicted			

**Overall Accuracy: 6/10 (60%)**

True			Overall Accuracy: 6/10 (60%)
	Cat	Dog	
Cat	4	1	4/5 (80%)
	3	2	2/5 (40%)
Cat		Dog	
Predicted			

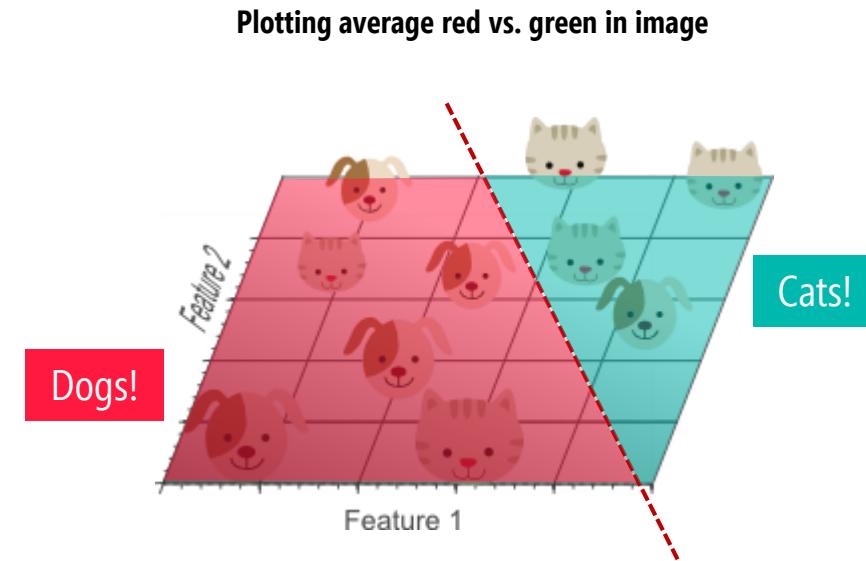
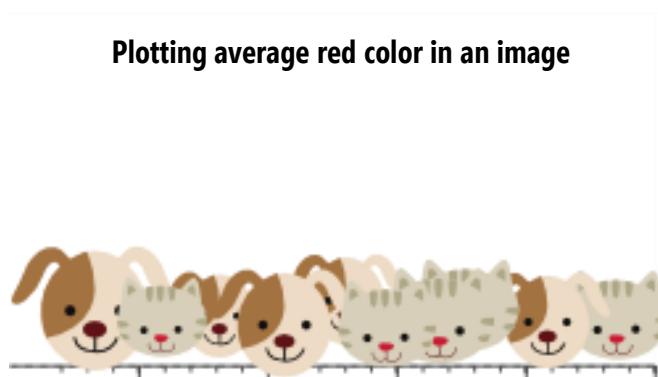
# EXERCISE: PLOT FEATURES IN 2D SPACE

Now in a new cell,  
start creating 2D  
scatter plots by  
combining your 1D  
features (assign one  
feature to x, another  
to y and graph!).



# SELECTING FEATURES: HIGHER DIMENSIONALITY

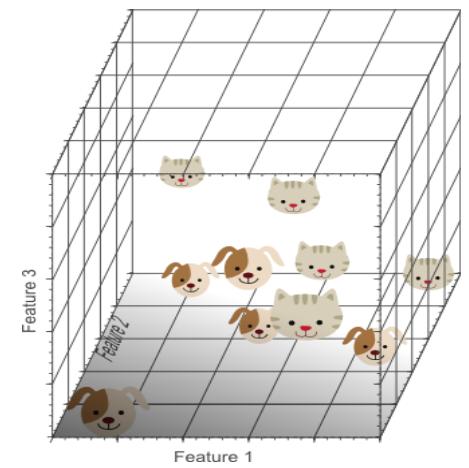
Let's look at another example: this time trying to separate cats and dogs using average pixel color in image.



## 1D Feature Vector

Plotting average red color in image.  
Notice how a single feature does not result in a perfect separation of our training data.

Plotting average red, green, blue in an image



## 2D Feature Vector

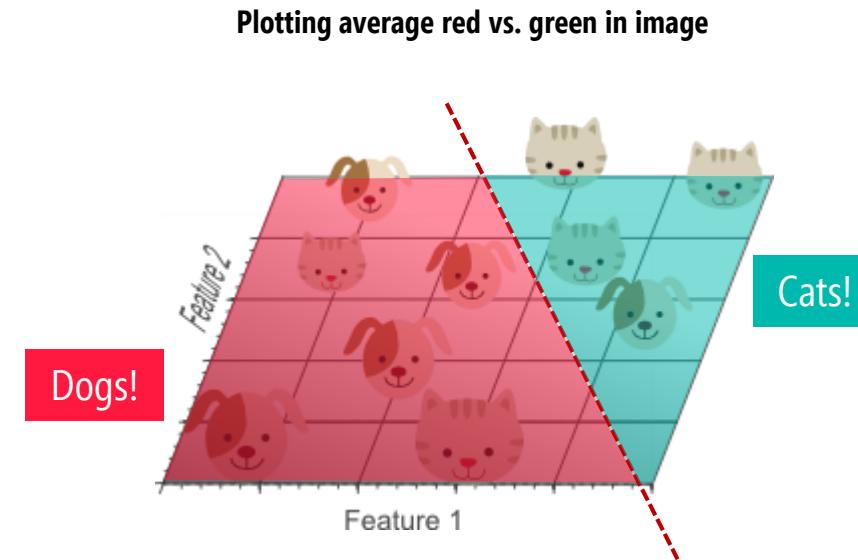
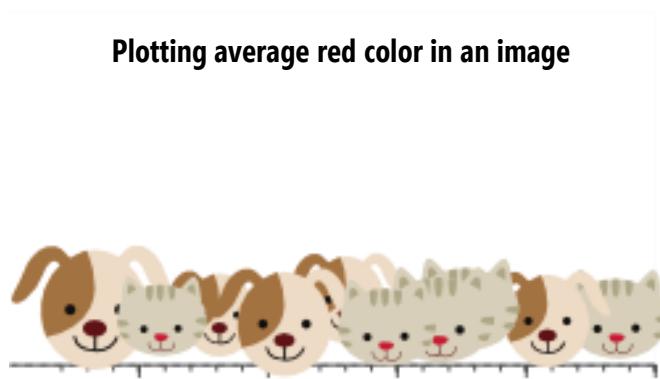
Plotting average red and green colors in image. Now we begin to see some more separation but not perfect.

## 3D Feature Vector

Using a 3D-input feature vector (R,G,B), we begin to see more separation. However, this is harder to see in the graph.

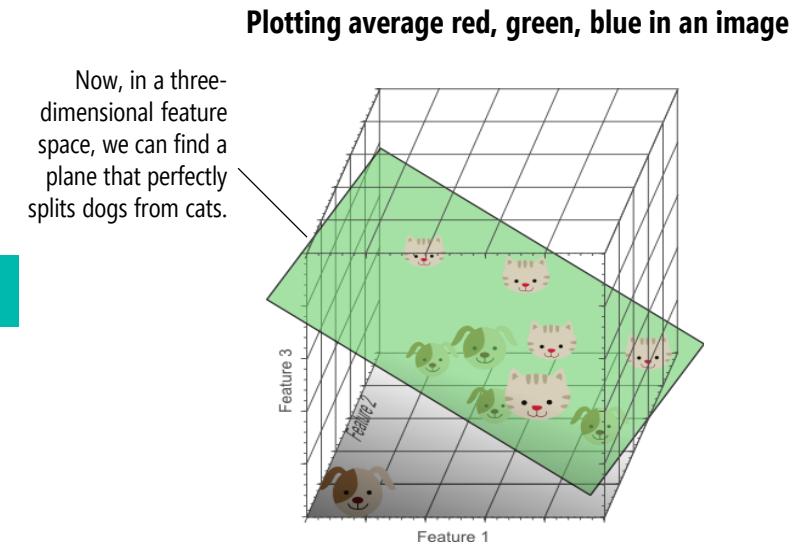
# SELECTING FEATURES: HIGHER DIMENSIONALITY

Let's look at another example: this time trying to separate cats and dogs using average pixel color in image.



## 1D Feature Vector

Plotting average red color in image.  
Notice how a single feature does not result in a perfect separation of our training data.



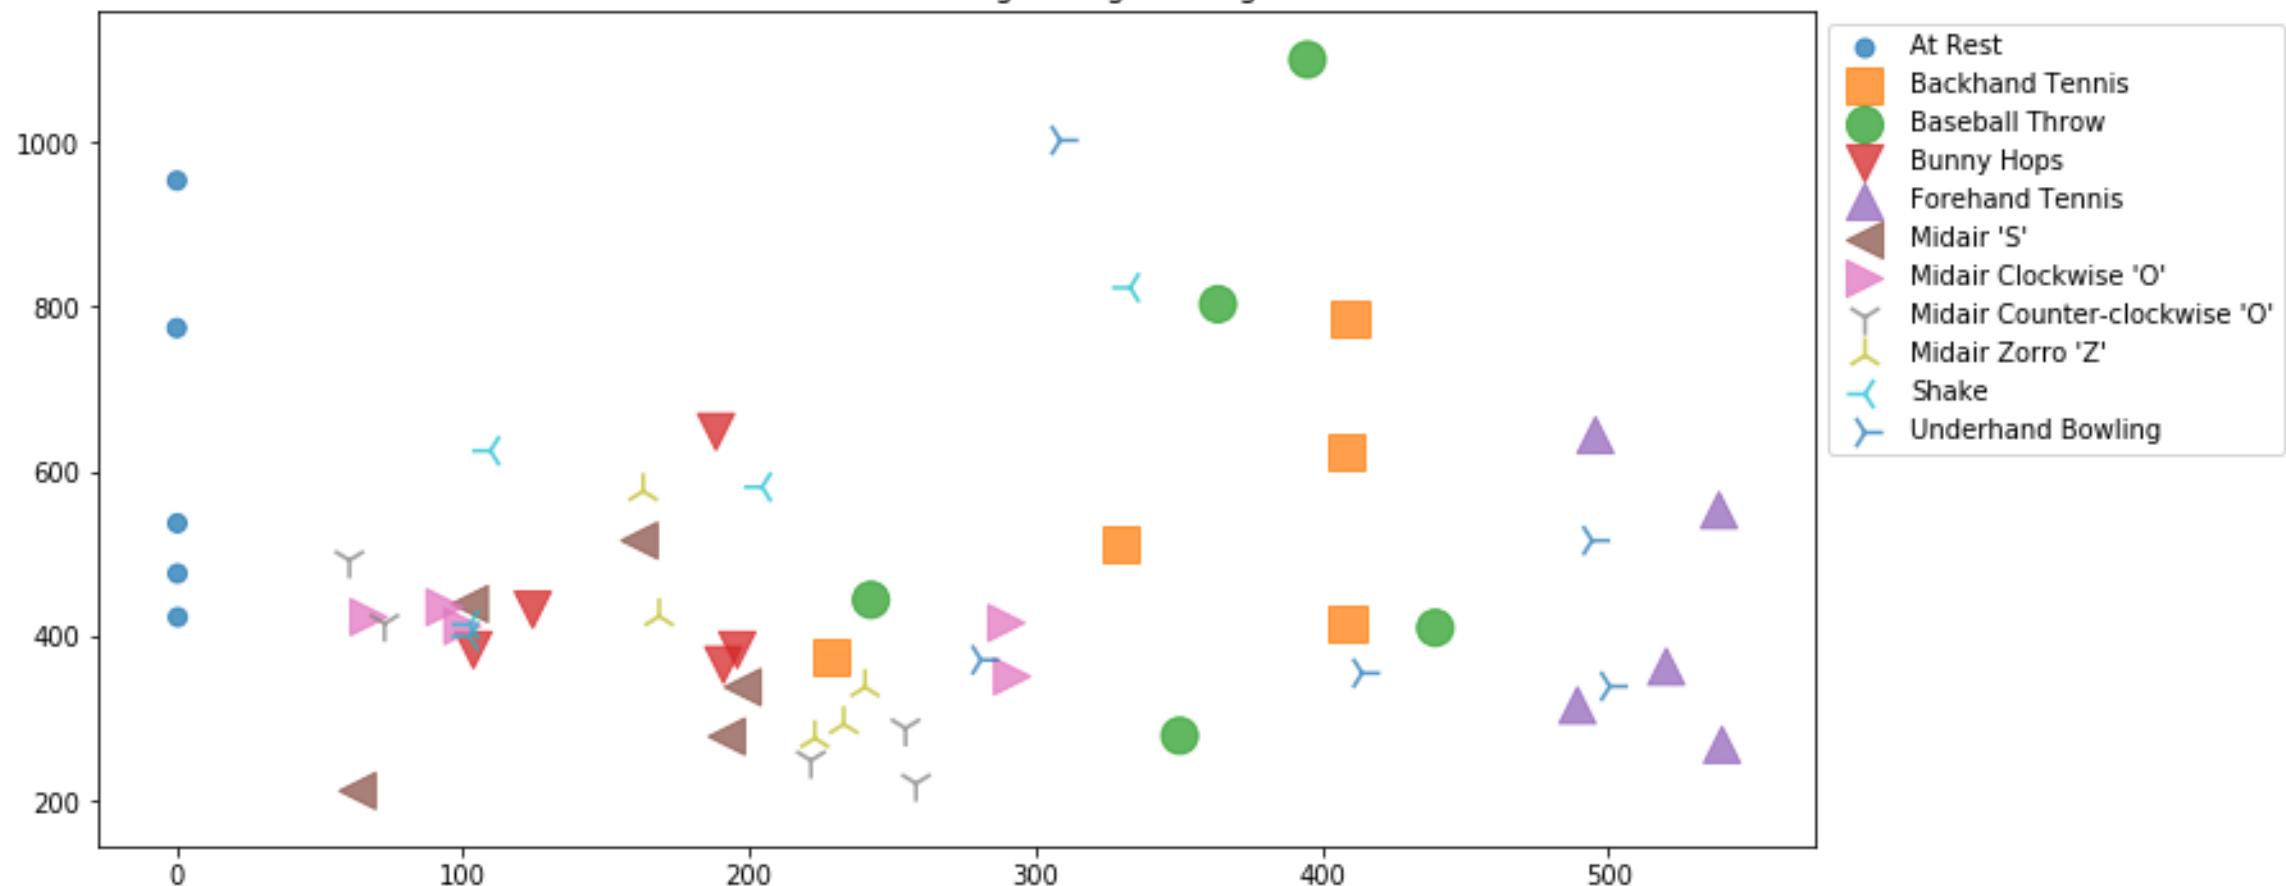
## 2D Feature Vector

Plotting average red and green colors in image. Now we begin to see some more separation but not perfect.

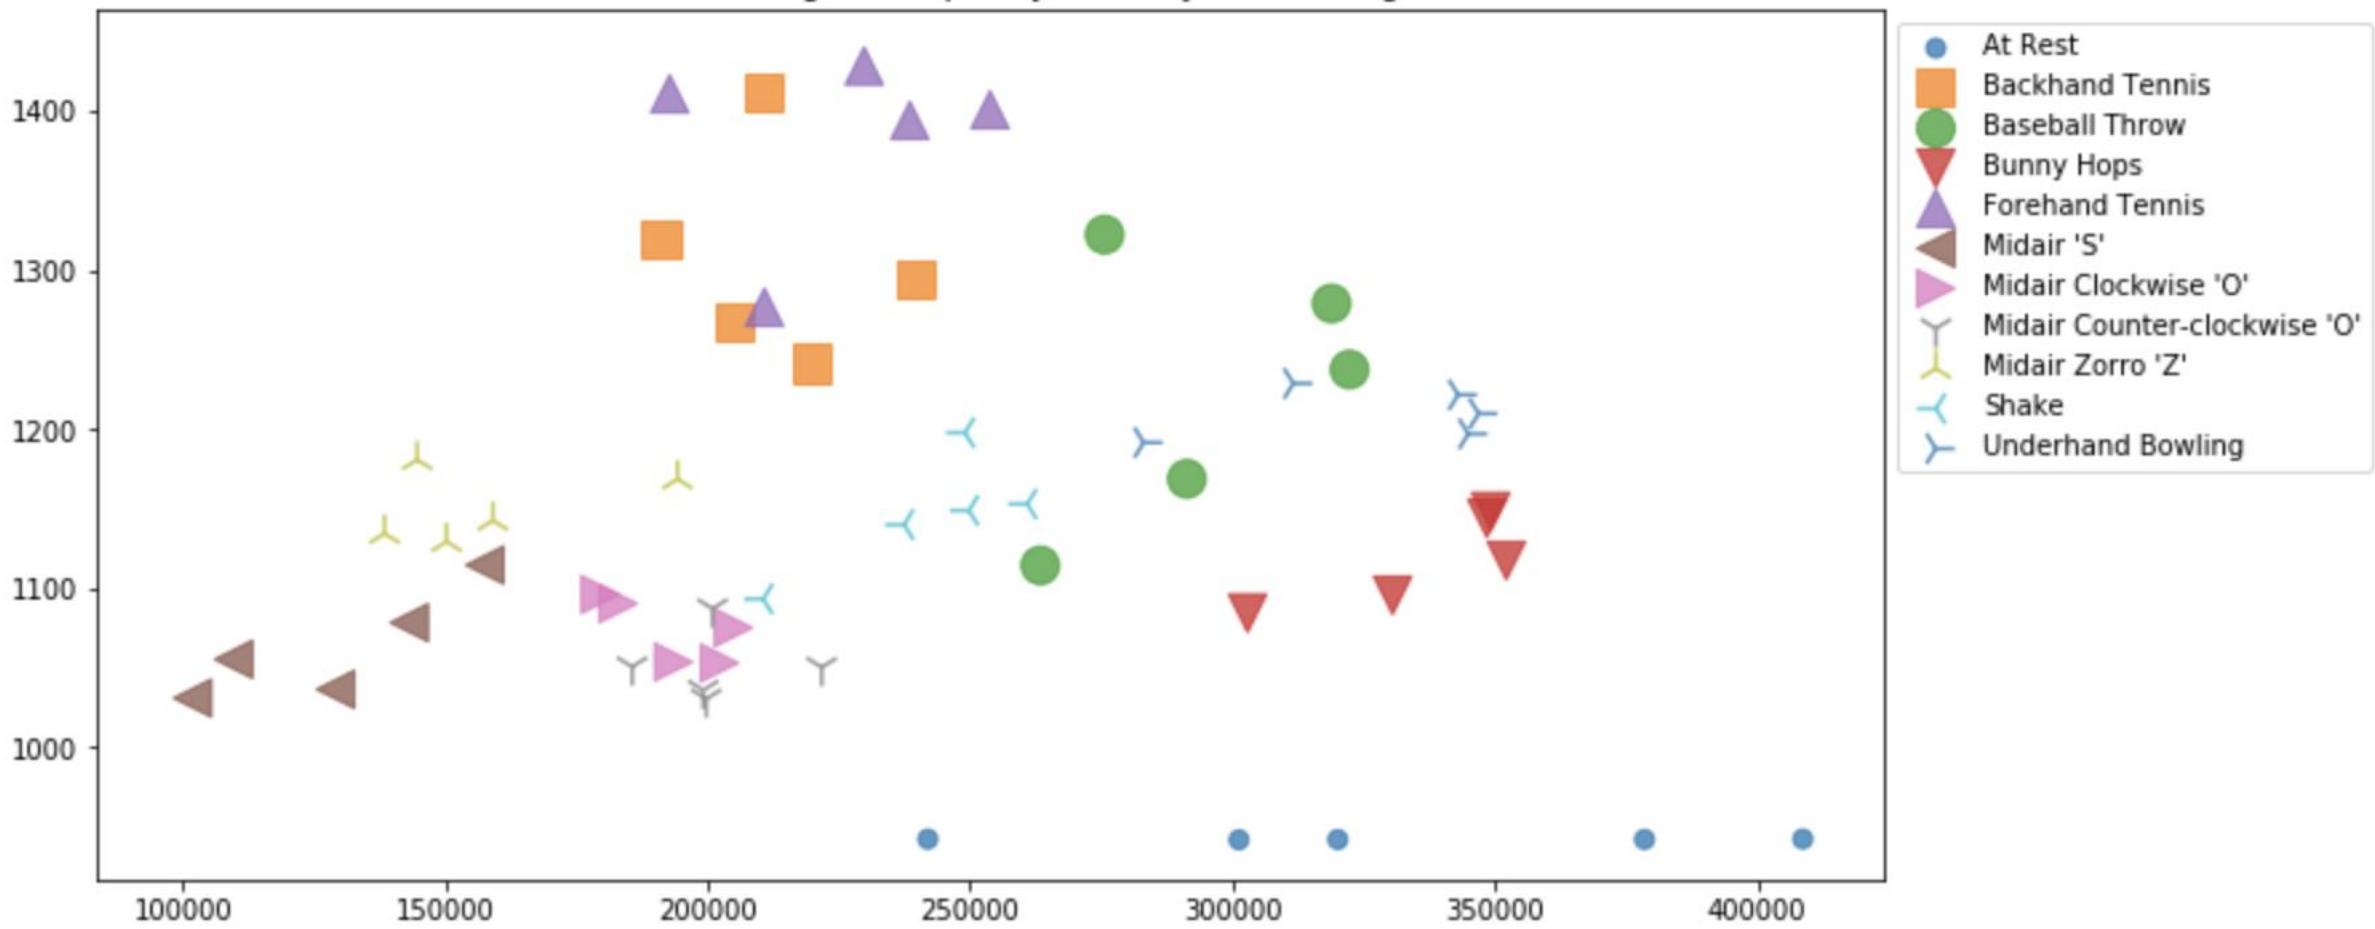
## 3D Feature Vector

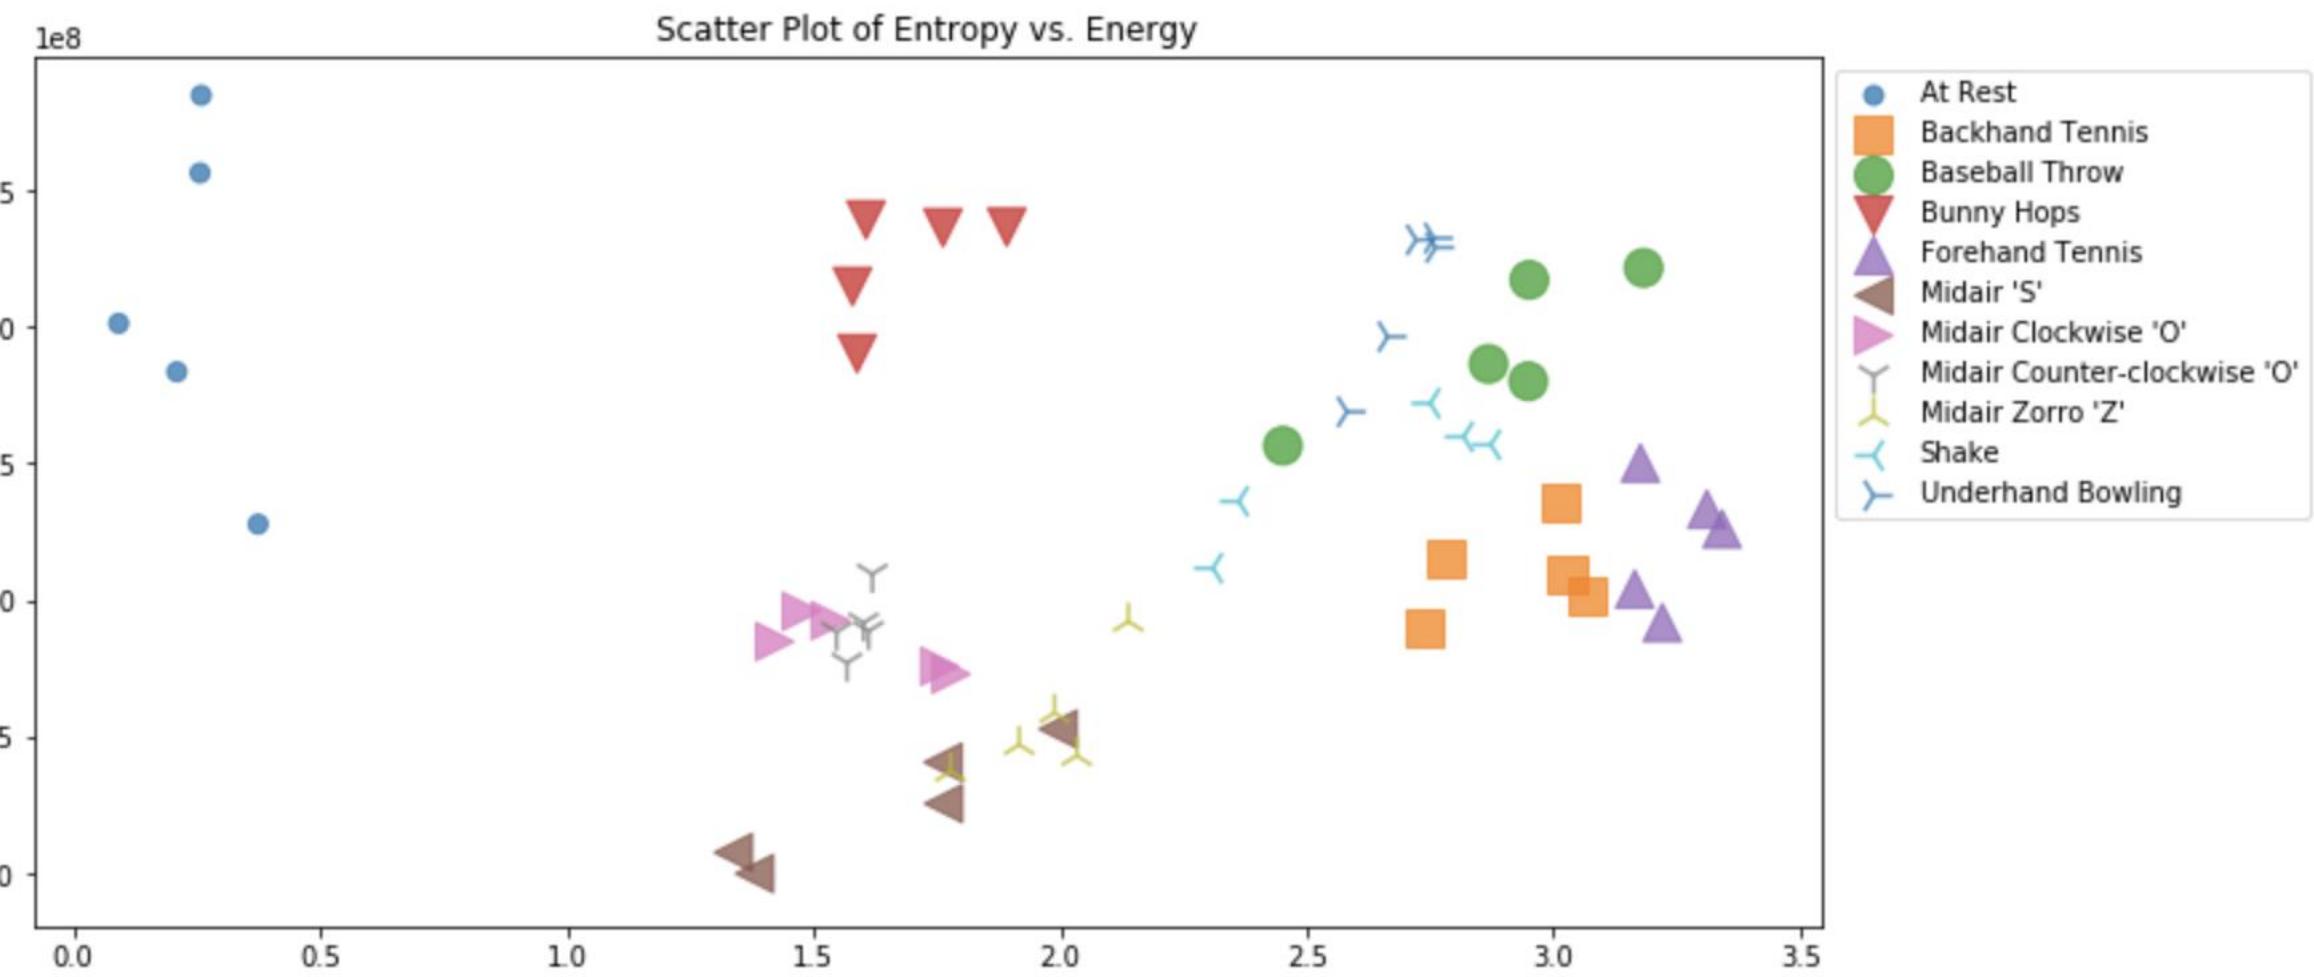
Using a 3D-input feature vector (R,G,B), we begin to see more separation. (Of course, plotting in 3D often requires interactivity to tilt, rotate, explore data)

2D Plot of Max Accel Mag vs. Signal Length

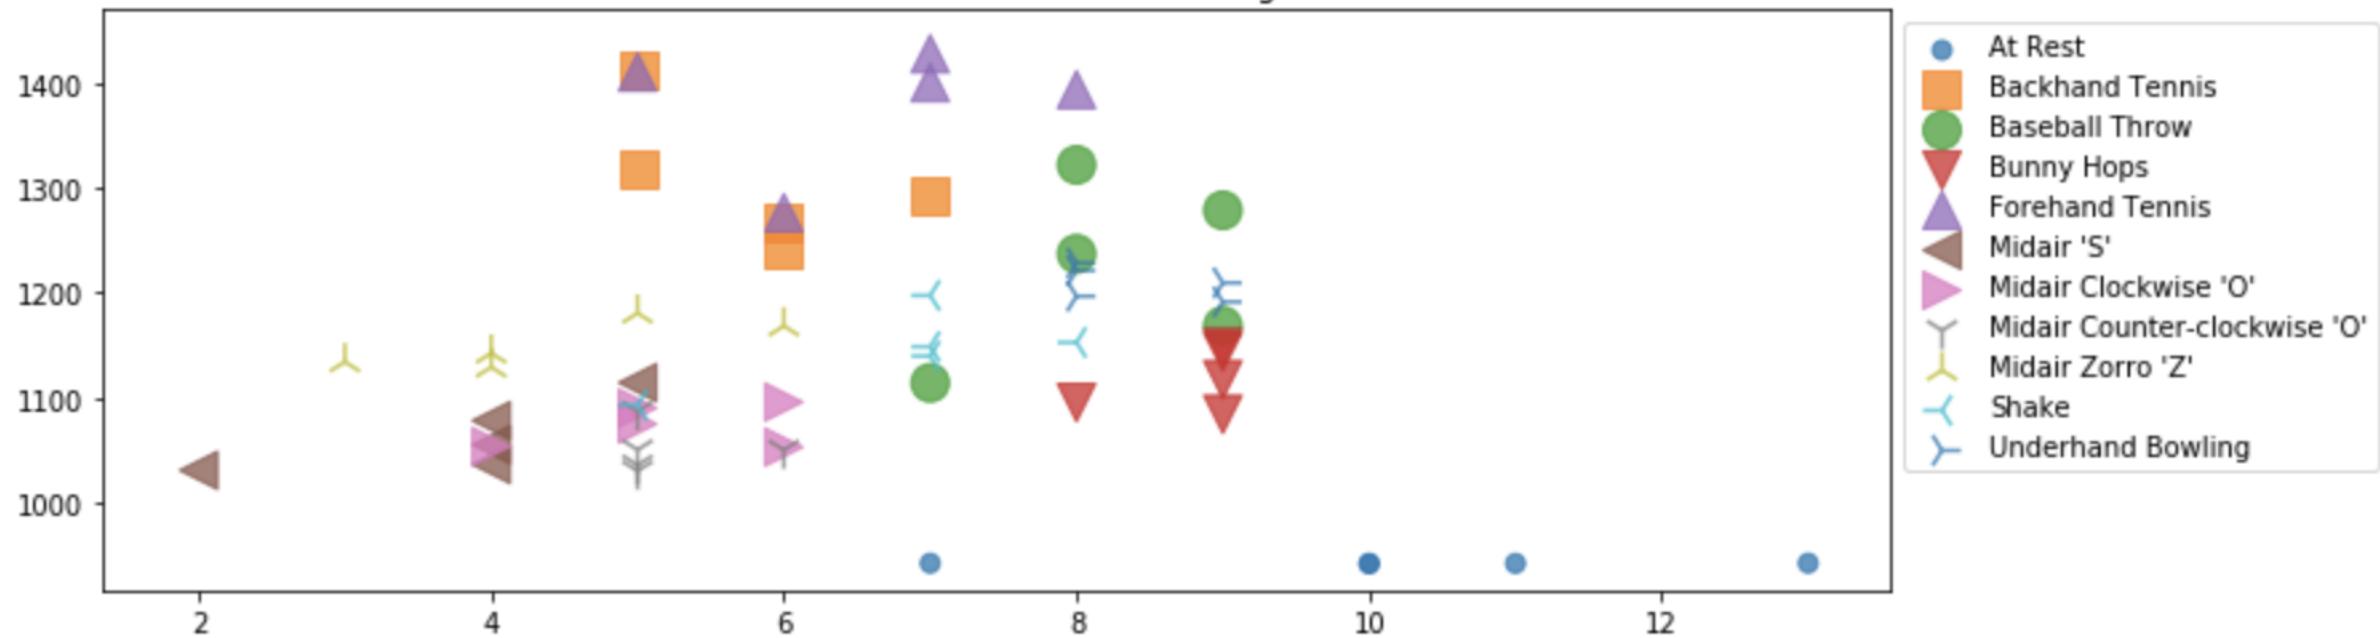


Scatter Plot of Strongest Frequency Intensity vs Max Magnitude





2D Plot of Num Peaks vs. Max Mag



# A3 REFLECTION + INTRO TO MODEL-BASED

---

CSE 599 Prototyping Interactive Systems | Lecture 14 | May 1

**Jon Froehlich** • Jasper O'Leary (TA)